

Dynamic Programming

Dynamic Programming

- Dynamic programming is an algorithmic technique for solving an optimization problem by breaking it down into simpler subproblems
- The optimal solution to the overall problem depends upon the optimal solution to its subproblems.
- Dynamic programming amounts to 1) breaking down a problem into simpler subproblems, and 2) storing the solution to each sub-problem so that each sub-problem is reused in similar problems.

- Example:



- $$c(n, k) = \frac{n!}{k!(n-k)!}$$
- $$c(n, k) = \begin{cases} c(n-1, k-1) + c(n-1, k), & \text{if } 0 < k < n \\ 1, & \text{if } k = 0 \text{ or } k = n \end{cases}$$

	0	1	2	3	4	k
0	1					
1	1	1				
2	1	2	1			
3	1	3	3	1		
4	1	4	6	4	1	
n						

- Dynamic Programming is a lot like divide and conquer approach
- The difference is results of a sub-problem are used in similar sub-problems.

Principle of Optimality

- Dynamic Programming is a very general solution method for problems which have two properties:
 - Optimal substructure
 - Principle of optimality applies
 - Optimal solution can be decomposed into subproblems
 - Overlapping subproblems
 - Subproblems recur many times
 - Solutions can be cached and reused
- Markov decision processes satisfy both properties
 - Bellman equation gives recursive decomposition
 - Value function stores and reuses solutions



Prediction & Control in Dynamic Programming

- Planning in an MDP
- For prediction: Prediction is to find the value function by evaluating a policy using the Bellman Expectation Equation.
 - Input: MDP $\langle S, A, P, R, \gamma \rangle$ and policy π
 - Output: **value function** v_π
- For control: This process involves optimizing the value function, we calculated during the prediction process.
 - Input: MDP $\langle S, A, P, R, \gamma \rangle$
 - Output: **optimal value function** v_* and **optimal policy** π_*
- Dynamic programming assumes full knowledge of the MDP
 - Assume **transition prob** and **rewards** are known.
 - Feasibility in real-world engineering applications is therefore limited

Policy Evaluation (Prediction)

- Prediction problem: given a policy π , compute the state-values(v) of the model using the policy

Method 1) Using matrix $\mathbf{v}_\pi = \mathbf{R}_\pi + \gamma \mathbf{P}_\pi \mathbf{v}_\pi \Rightarrow \mathbf{v}_\pi = (\mathbf{I} - \gamma \mathbf{P}_\pi)^{-1} \mathbf{R}_\pi$
(from MDP chapter)

Method 2) Iterative approach: Iterative application of Bellman Expectation Equation

$$v_1 \rightarrow v_2 \rightarrow \cdots \rightarrow v_\pi$$

Policy Evaluation (Prediction)

Matrix Method:

- Use Bellman expectation equation
- Recap: Bellman Expectation Equation:

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) \left(R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_{\pi}(s') \right) = \underbrace{\sum_{a \in A} \pi(a|s) R_s^a}_{R_s^{\pi}} + \gamma \sum_{s' \in S} \underbrace{\sum_{a \in A} \pi(a|s) P_{ss'}^a}_{P_{ss'}^{\pi}} v_{\pi}(s')$$

- For every state s , define $v_{\pi}(s)$

$$\begin{bmatrix} v_{\pi}(1) \\ \vdots \\ v_{\pi}(n) \end{bmatrix} = \begin{bmatrix} R_1^{\pi} \\ \vdots \\ R_n^{\pi} \end{bmatrix} + \gamma \begin{bmatrix} P_{11}^{\pi} & \cdots & P_{1n}^{\pi} \\ \vdots & \vdots & \vdots \\ P_{n1}^{\pi} & \cdots & P_{nn}^{\pi} \end{bmatrix} \begin{bmatrix} v_{\pi}(1) \\ \vdots \\ v_{\pi}(n) \end{bmatrix}$$

$$\mathbf{v}_{\pi} = \mathbf{R}_{\pi} + \gamma \mathbf{P}_{\pi} \mathbf{v}_{\pi} \Rightarrow \mathbf{v}_{\pi} = (\mathbf{I} - \gamma \mathbf{P}_{\pi})^{-1} \mathbf{R}_{\pi}$$

Policy Evaluation (Prediction)

Iterative Method:


- When number of states is too big or continuous, matrix method is too complex or impossible
- During one iteration, the *old* value of s is replaced with a *new* value from the old values of the successor state s'
 - Update $v_{k+1}(s)$ from $v_k(s)$
- Algorithm
 - At each iteration $k+1$
 - For all states $s \in S$
 - Update $v_{k+1}(s)$ from $v_k(s')$ where s' is a successor state of s

$$v_{k+1}(s) = \sum_{a \in A} \pi(a|s) \left(R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_k(s') \right)$$

Iterative Policy Evaluation

- Updating estimates v_{k+1} on the basis of other estimates v_k is often called **bootstrapping**.
- This leads to synchronous, full backups of the entire state space.
- Using synchronous backups, **(synchronous: update all v values simultaneously)**
- Convergence to v_π is guaranteed

Evaluating a Random Policy in the Small Gridworld



- Undiscounted episodic MDP ($\gamma=1$)
- Nonterminal states $1, \dots, 14$
- One terminal state (shown twice as shaded squares)
- Actions leading out of the grid leave state unchanged
- Reward is -1 until the terminal state is reached
- Agent follows uniform random policy

$$\pi(n | \cdot) = \pi(e | \cdot) = \pi(s | \cdot) = \pi(w | \cdot) = 0.25$$

Evaluating a Random Policy in the Small Gridworld

- Assume deterministic transition

$$\sum_{a \in A} P_{ss'}^a v_k(s') = v_k(s')$$

- Iterative Method

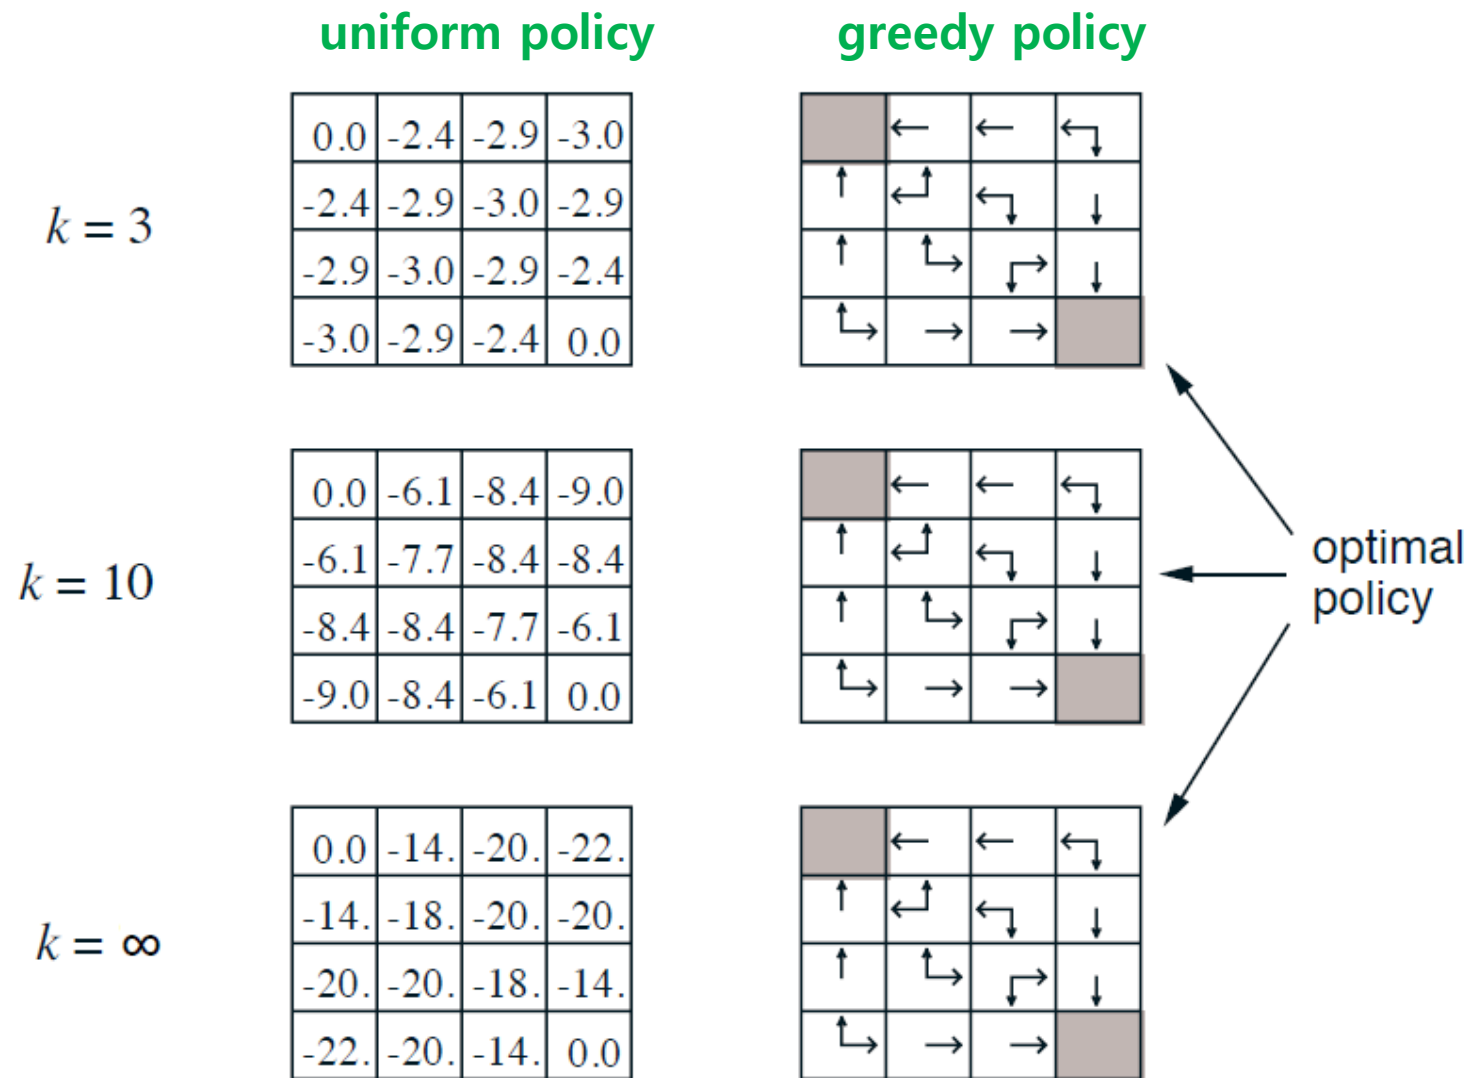
$$v_{k+1}(s) = \sum_{a \in A} \pi(a|s) \left(R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_k(s') \right) = \sum_{a \in A} \pi(a|s) (-1 + v_k(s'))$$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$$v_{k+1}(s_1) = -1 + 0.25 * (0 + (-6.1) + (-8.4) + (-7.7)) = -6.55$$

$$v_{k+1}(s_6) = -1 + 0.25 * ((-7.7) + (-8.4) + (-8.4) + (-7.7)) = -9.05$$

Iterative Policy Evaluation in Small Gridworld



R. Sutton and A. Barto, Reinforcement Learning: An Introduction, 2018

Control Problem

- We have seen how to evaluate a given policy
- However, **our main goal is to find an optimal policy.** How?
- Control problem:
- **Policy Iteration**
 - Policy Evaluation: given a policy, evaluate it (compute state-values)
 - Policy Improvement: improve policy
- **Value Iteration**
 - Compute state values (no policy is given)
 - At the end, derive optimal policy
 - No need of policy evaluation

Policy Iteration

- Policy Iteration

- 1) Given an arbitrary policy π

- 2) **Evaluation:** evaluate the policy π (= policy prediction)

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) (R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_{\pi}(s'))$$


- 3) **Improvement:** improve the policy by acting greedily with respect to v_{π}

$$\pi' = \text{greedy}(v_{\pi})$$

- 4) We iterate these two processes 2)-3) until it converges


- Each policy evaluation step is fully executed, i.e. for each policy π_i an exact estimate of v_{π_i} is computed either by iterative method or by any other method

Policy Iteration



Policy evaluation Estimate v_π
Iterative policy evaluation

Policy improvement Generate $\pi' \geq \pi$
Greedy policy improvement



R. Sutton and A. Barto, Reinforcement Learning: An Introduction, 2018

Greedy Policy Improvement

- Consider a deterministic policy, $a = \pi(s)$ (Initial policy may be stochastic)
- Once policy evaluation is done, we can *improve* the policy by acting greedily
 - Select the best action according to $q(s_k, a_k)$ in every state

$$\pi'(s) = \operatorname{argmax}_{a \in A} q_{\pi}(s, a)$$

- π' becomes a deterministic policy
- Is taking the greedy action $\pi'(s)$ is better than just following our policy π ?
 - This improves the value from any state s over **one step**

$$q_{\pi}(s, \pi'(s)) = \max_{a \in A} q_{\pi}(s, a) \geq q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$

- **$\pi(s)$ is a deterministic policy**
- It therefore improves the value function, $v_{\pi'}(s) \geq v_{\pi}(s)$

Greedy Policy Improvement

- **Policy Improvement Theorem:** Greedy policy improvement improves the value function, $v_{\pi'}(s) \geq v_{\pi}(s)$

Proof: $v_{\pi}(s) \leq q_{\pi}(s, \pi'(s))$ (from previous page)

$$\begin{aligned} &= E_{\pi'}[R_{t+1} + \gamma v_{\pi}(s_{t+1}) | S_t = s] \\ &\leq E_{\pi'}[R_{t+1} + \gamma q_{\pi}(s_{t+1}, \pi'(s_{t+1})) | S_t = s] \\ &\leq E_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 q_{\pi}(s_{t+2}, \pi'(s_{t+2})) | S_t = s] \\ &\leq E_{\pi'}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s] = v_{\pi'}(s) \end{aligned}$$

Policy Improvement

- If improvements stop means, no changes in q or v values
 - Thus reached their maximum
 - Does it reach optimum policy?

$$\underbrace{q_{\pi}(s, \pi'(s))}_{\text{definition of } \pi'(s)} = \max_{a \in A} q_{\pi}(s, a) \stackrel{\text{no change}}{=} \underbrace{q_{\pi}(s, \pi(s))}_{\text{definition of } v_{\pi}(s)} = v_{\pi}(s) \text{ (policy doesn't change)}$$

- If the Bellman optimality equation has been satisfied

$$v_{\pi}(s) = \max_{a \in A} q_{\pi}(s, a)$$

- then $v_{\pi}(s) = v_{*}(s)$ for all $s \in S$
- So π is an optimal policy
- Policy improvement theorem guarantees finding optimal policies in finite MDPs



Modified Policy Iteration

- Does policy evaluation need to converge to v_π ?
- Or should we introduce a stopping condition
 - e.g. ϵ -convergence of value function
- Or simply stop after k iterations of iterative policy evaluation?
- For example, in the small grid-world $k=3$ was sufficient to achieve optimal policy
- In value iteration, for example, only a *single iteration* of policy evaluation is performed in between each policy improvement.
- As long as both processes continue to update all states, the ultimate result is typically the same
 - convergence to the optimal value function and an optimal policy.

Generalized Policy Iteration

- Generalized policy iteration(GPI) refers to all algorithms based on policy iteration
- Almost all reinforcement learning methods are well described as GPI.
- The policy always being improved with respect to the value function and the value function always being driven toward the value function for the policy
- If both the evaluation process and the improvement process stabilize, then the value function and policy must be optimal.
- Guaranteed to converge to the optimal policy, provided PE and PI are executed enough times.

Generalized Policy Iteration



Policy evaluation Estimate v_π

Any policy evaluation algorithm


Policy improvement Generate $\pi' \geq \pi$

Any policy improvement algorithm

R. Sutton and A. Barto, Reinforcement Learning: An Introduction, 2018

Principle of Optimality

- Any optimal policy can be subdivided into two components:
 - An optimal 1st action A
 - Followed by an optimal policy from successor state S'



Theorem (Principle of Optimality)

A policy $\pi(a|s)$ achieves the optimal value from state s , $v_\pi(s) = v_*(s)$, if and only if

- For any state s' reachable from s
- π achieves the optimal value from state s' , $v_\pi(s') = v_*(s')$

$$v_*(s) = \max_a (R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_*(s'))$$

Value Iteration

- Policy iteration involves full policy evaluation steps between policy improvements.
- In large state-space MDPs, the full policy evaluation may be numerically very costly.
- Using a limited number of iterative policy evaluation steps and then apply policy improvement may speed up the entire DP process.
- **Value iteration**: the special case for Policy Iteration. The process of policy evaluation is stopped after one step.

Value Iteration

- Iterative application of Bellman optimality equation ($v_1 \rightarrow v_2 \rightarrow \dots v_*$)
- Set $k=1$; Initialize $V_0(s) = 0$ for all states s
 - At each iteration $k+1$ [until converge]
 - For all states $s \in S$


$$v_{k+1}(s) = \underbrace{\max_{a \in A}}_{\text{Policy improvement}} \left(\underbrace{R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_k(s')}_{\text{Policy evaluation}} \right)$$

- Output: deterministic policy π such that

$$\pi_*(s) = \operatorname{argmax}_{a \in A} \left(R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_{k+1}(s') \right)$$

- Converges to v_*
- Unlike policy iteration, there is no explicit policy
- Intermediate value functions may not correspond to any policy

Value Iteration



Bellman Optimality Equation

$$v_*(s) = \max_a (R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a v_*(s'))$$

Principle of Optimality

$$v_{k+1}(s) = \max_{a \in A} \left(R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a v_k(s') \right)$$

$$\mathbf{v}_{k+1} = \max_{a \in A} (\mathbf{R}^a + \gamma \mathbf{P}^a \mathbf{v}_k)$$

Policy improvement + policy evaluation

Summarizing DP Algorithms

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- All DP Algorithms are based on state-value function $v_{\pi}(s)$ or $v_{*}(s)$
 - Complexity $O(mn^2)$ per iteration for m actions and n states
 - Evaluate all n^2 state transitions while considering up to m actions per state.
- Could also apply to action-value function $q_{\pi}(s, a)$ or $q_{*}(s, a)$
 - Complexity $O(mn^3)$ per iteration
 - There are up to nm action-values which require n^2 state transition evaluations each.

Asynchronous Dynamic Programming

- DP algorithms considered so far used **synchronous backups**:
 - In one iteration the entire state space is updated.
 - Computational expensive for large MDPs.
 - Some state-values or policy parts may converge faster than other but are updated as often as slowly converging states.
- **Asynchronous backups** update states individually in an (arbitrary) order:
 - Some states may be updated more frequently than others.
 - Choose smart order to achieve faster overall convergence rate.
 - Overall algorithms converges if all states are still visited to some extent
- Asynchronous DP backs up states individually, in any order
- For each selected state, apply the appropriate backup
- Can significantly reduce computation
- Guaranteed to converge if all states continue to be selected

Curse of Dimensionality

- DP uses full-width backups:
 - For each state update, every successor state and action is considered.
 - While utilizing full knowledge of the MDP structure.
- Hence, DP is can be effective up to medium-sized MDPs
- Also, we must have full knowledge of MDP structure
- For large problems DP suffers from the curse of dimensionality:
 - Number of finite states n grows exponentially with the number of state variables.
 - Also: if continuous variables need quantization typically a large number of states results.