

Model-Free Control

Model-Free RL

- Model-free prediction: how good is a specific policy?
 - Given no access to the decision process model parameters
 - Instead have to estimate from data/experience
- Model-free control: how can we learn a good policy?

Target Policy vs Behavior Policy

- All learning control methods face a dilemma: They learn action values for optimal behavior, but need to behave non-optimally to explore new actions.
- A more straightforward approach is to use two policies: The policy being learned about is called the **target policy**, and the policy used to generate behavior is called the **behavior policy**.
 - Target policy: It is the policy that an agent is trying to evaluate or improve. i.e agent is learning value function for this policy.
 - Behavior policy: It is the policy that is being used by an agent for action select. i.e agent follows this policy to interact with the environment.
- The target policy is typically the *deterministic* greedy policy with respect to the current estimate of the action-value function.
- The behavior policy remains *stochastic* and more exploratory, for example, an ϵ -greedy policy.


On-Policy vs Off-Policy

- How to learn about the optimal policy while behaving according to an exploratory policy?
- On-policy
 - target policy = behavior policy
 - learn to estimate and evaluate a policy from experience obtained from following that policy
 - learns action values not for the optimal policy, but for a near-optimal policy that still explores.
 - generally simpler and are considered first
 - e.g. Sarsa


On-Policy vs Off-Policy

- Off-policy
 - target policy \neq behavior policy
 - learn to estimate and evaluate a policy using experience gathered from following a **different** policy
 - more powerful and general
 - can be applied to learn from data generated by a conventional non-learning controller, or from a human expert.
 - re-use experience generated from old policies (π_0, π_1, \dots) .
 - are often of greater variance and are slower to converge
 - e.g. Q learning

Recap: Generalized Policy Iteration




- Policy evaluation Estimate V_π
 - e.g. Iterative policy evaluation
- Policy improvement Generate $\pi' \geq \pi$
 - e.g. Greedy policy improvement



R. Sutton and A. Barto, Reinforcement Learning: An Introduction, 2018

Generalized Policy Iteration With Monte-Carlo Evaluation



- We already learned **model-free policy evaluation** method(e.g. MC)
- Also learned **greedy policy improvement**
- Therefore, how about this GPI?
 - Policy evaluation: **Monte-Carlo** policy evaluation, $V = v_\pi$?
 - Policy improvement: **Greedy** policy improvement?

Model-Free Policy Iteration Using Action-Value Function


- Greedy policy improvement over $V(s)$ requires model of MDP

$$\pi'(s) = \operatorname{argmax}_{a \in A} R_s^a + P_{ss'}^a V(s')$$

- Greedy policy improvement over $Q(s,a)$ is model-free
 - We **do not** know $P_{ss'}^a$, (assuming stochastic) and R_s^a
 - We need **action-value** function for **model-free** control

$$\pi'(s) = \operatorname{argmax}_{a \in A} Q(s, a)$$

Generalized Policy Iteration With Monte-Carlo Evaluation



- Now use Q instead of V
- Policy evaluation: Monte-Carlo policy evaluation, $Q = q_\pi$?
- Policy improvement: Greedy policy improvement?
- Does it work now?

Problem of Greedy Policy Improvement

- Initialize policy π
- Repeat:
 - **Policy evaluation**: compute $Q_{\pi_i}(s, a) \forall s, a$
 - **Policy improvement**: greedy update π given Q_{π_i}
$$\pi'(s) = \operatorname{argmax}_{a \in A} Q(s, a)$$
- **Greedy** method for policy improvement -> deterministic policy
- If π is deterministic, can't compute $Q(s, a)$ for any $a \neq \pi(s)$
 - May need to modify policy evaluation
 - Policy improvement is using an *estimated* Q
- How to interleave policy evaluation and improvement?

Example of Greedy Action Selection



"Behind one door is tenure - behind the other is flipping burgers at McDonald's."

<https://www.davidsilver.uk/teaching/>

- There are two doors in front of you.
- You open the left door and get reward 0
 $V(\text{left}) = 0$
- You open the right door and get reward +1
 $V(\text{right}) = +1$
- You open the right door and get reward +3
 $V(\text{right}) = +2$
- You open the right door and get reward +2
 $V(\text{right}) = +2$
- ...
- Are you sure you've chosen the best door?

ϵ -greedy Exploration

- If π is a deterministic policy, then in following π one will observe returns only for one of the actions from each state
 - many state–action pairs may never be visited.
- Need to try all (s, a) pairs but then follow π
- Want to ensure resulting estimate Q_π is good enough so that policy improvement is a monotonic operator
- **ϵ -greedy exploration:** Simple idea to balance exploration and exploitation
- All m actions are tried with non-zero probability
 - With probability $1 - \epsilon$, choose the greedy action
 - With probability ϵ , choose an action at random

$$\pi(a|s) = \begin{cases} \frac{\epsilon}{m} + 1 - \epsilon, & \text{if } a^* = \operatorname{argmax}_{a \in A} Q(s, a) \\ \frac{\epsilon}{m}, & \text{otherwise} \end{cases}$$

ϵ -greedy Policy Improvement

- **Theorem:** For any ϵ -greedy policy π , the ϵ -greedy policy π' with respect to q_π is an improvement, $v_{\pi'}(s) \geq v_\pi(s)$

$$q_\pi(s, \pi'(s)) = \sum_a \pi'(a|s) q_\pi(s, a) = \frac{\epsilon}{m} \sum_a q_\pi(s, a) + (1 - \epsilon) \max_a q_\pi(s, a) \frac{1 - \epsilon}{1 - \epsilon}$$

- Since


$$\sum_a (\pi(a|s) - \epsilon/m) = \sum_{a \neq a_*} (\pi(a|s) - \epsilon/m) + \sum_{a=a_*} (\pi(a|s) - \epsilon/m) = (0) + (1 - \epsilon + \frac{\epsilon}{m} - \frac{\epsilon}{m}) = 1 - \epsilon$$

- Therefore,

$$\begin{aligned} q_\pi(s, \pi'(s)) &= \frac{\epsilon}{m} \sum_a q_\pi(s, a) + (1 - \epsilon) \max_a q_\pi(s, a) \frac{\sum_a (\pi(a|s) - \epsilon/m)}{1 - \epsilon} \\ &\geq \frac{\epsilon}{m} \sum_a q_\pi(s, a) + (1 - \epsilon) \sum_a \frac{\pi(a|s) - \epsilon/m}{1 - \epsilon} q_\pi(s, a) \\ &= \sum_a \pi(a|s) q_\pi(s, a) = v_\pi(s) \end{aligned}$$


- Since $v_{\pi'}(s) \geq q_\pi(s, \pi'(s))$, $v_{\pi'}(s) \geq v_\pi(s)$

Monte-Carlo Policy Iteration



- Now use ϵ -greedy policy improvement instead of greedy improvement
 - Policy evaluation: Monte-Carlo policy evaluation, $Q = q_\pi$
 - (Compute q values until it converge)
 - Policy improvement: ϵ -greedy policy improvement
- This works.

Monte-Carlo Control



- We can do better (more efficiently).
- For every episode:
 - Policy evaluation: Monte-Carlo policy evaluation, $Q \approx q_\pi$
 - No need to compute entire q values
 - Policy improvement: ϵ -greedy policy improvement

Greedy in the Limit with Infinite Exploration (GLIE)

- Greedy in the limit with infinite exploration (GLIE)
- A learning policy π is called GLIE if it satisfies the following two properties:
 - All state-action pairs are explored infinitely many times

$$\lim_{k \rightarrow \infty} N_k(s, a) = \infty$$

- The policy converges on a greedy policy

$$\lim_{k \rightarrow \infty} \pi_k(a|s) = 1(a = \operatorname{argmax}_{a'} Q_k(s, a'))$$

- For example, ϵ -greedy is GLIE if ϵ reduces to zero at $\epsilon_k = \frac{1}{k}$

GLIE Monte Carlo Control

- Sample k th episode using $\pi: \{S_1, A_1, R_2, \dots S_T\} \sim \pi$
- For each state S_t and action A_t in the episode,

$$\begin{aligned} N(S_t, A_t) &\leftarrow N(S_t, A_t) + 1 \\ Q(S_t, A_t) &\leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)} (G_t - Q(S_t, A_t)) \end{aligned}$$

- Improve policy based on new action-value function

$$\begin{aligned} \epsilon &\leftarrow 1/k \\ \pi &\leftarrow \epsilon\text{-greedy}(Q) \end{aligned}$$

Theorem: GLIE Monte-Carlo control converges to the optimal action-value function, $Q(s, a) \rightarrow Q_*(s, a)$


MC vs. TD Control

- Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)
 - Lower variance
 - Incomplete sequences
 - Online
- Natural idea: use TD instead of MC in our control loop
 - Apply TD to $Q(S,A)$
 - Use ϵ -greedy policy improvement
 - Update every time-step

Model-Free Policy Iteration with TD Methods


- Use temporal difference methods for policy evaluation step
- Initialize policy π
- Repeat:
 - Policy evaluation: compute Q_π using temporal difference(TD)
 - Policy improvement: set π to ϵ -greedy (Q_π) (same as Monte Carlo policy improvement)
- First consider SARSA, which is an on-policy algorithm.

On-Policy Control With Sarsa



- Instead of MC, use **Sarsa** & update q at **every step** (instead of every episode)
- For every **time-step**:
 - Policy evaluation: **Sarsa**, $Q \approx q_{\pi}$
 - Policy improvement: ϵ -greedy policy improvement

Updating Action-Value Functions with Sarsa




$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$

Sarsa Algorithm for On-Policy Control

- For every **time-step**:
 - Policy evaluation: **Sarsa**, $Q \approx q_\pi$
 - Policy improvement: ϵ -greedy policy improvement

```
1: Set initial  $\epsilon$ -greedy policy  $\pi$ ,  $t = 0$ , initial state  $s_t = s_0$ 
2: Take  $a_t \sim \pi(s_t)$  // Sample action from policy
3: Observe  $(r_t, s_{t+1})$ 
4: loop
5:   Take action  $a_{t+1} \sim \pi(s_{t+1})$ 
6:   Observe  $(r_{t+1}, s_{t+2})$ 
7:    $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$ 
8:    $\pi(s_t) = \arg \max_a Q(s_t, a)$  w.prob  $1 - \epsilon$ , else random
9:    $t = t + 1$ 
10: end loop
```



ϵ -greedy policy improvement

policy evaluation (Sarsa)

Convergence of Sarsa Algorithm




- Sarsa converges to the optimal action-value function, $Q(s, a) \rightarrow Q_*(s, a)$, under the following conditions:

- 1) Policy sequence $\pi_t(a|s)$ satisfies the condition of GLIE
- 2) The step-sizes α_t satisfy Robbins-Monro sequence such that

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$
$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

- For example, $\alpha_t = 1/t$ satisfies the above condition

Windy Gridworld




R. Sutton and A. Barto, Reinforcement Learning: An Introduction, 2018

- Reward = -1 per time-step until reaching goal
- Undiscounted
- Until $t=2000$, it couldn't finish episodes, but after that number of finished episodes grows exponentially

Expected Sarsa

- Expected Sarsa is very similar to Sarsa.
- However, instead of state-action values sampled using our current policy, it computes the expected value over all future state-action pairs
- The update-step is now defined as:

$$q(s_t, a_t) \leftarrow q(s_t, a_t) + \alpha \left(r_{t+1} + \underbrace{\gamma \sum_{a'} \pi(a'|s_{t+1}) q(s_{t+1}, a')}_{\text{average q value over action } a'} - q(s_t, a_t) \right)$$



- Expected Sarsa is an off-policy algorithm

Off-policy Learning

- Evaluate target policy $\pi(a|s)$ to compute $v_\pi(s)$ or $q_\pi(s, a)$ while following behavior policy $b(a|s)$

$$(S_1, A_1, R_1, \dots S_T) \sim b$$

- Why is this important?
 - Learn from observing humans or other agents
 - Re-use experience generated from old policies $\pi_1, \pi_2, \dots \pi_{t-1}$
 - Learn about optimal policy while following exploratory policy

Importance Sampling

- Some off-policy methods utilize importance sampling, a general technique for estimating expected values under one distribution given samples from another
- Estimate the expectation of a different distribution

$$E_{X \sim P}[f(X)] = \sum P(X)f(X) = \sum Q(X) \frac{P(X)}{Q(X)} f(X) = E_{X \sim Q} \left[\frac{P(X)}{Q(X)} f(X) \right]$$

Importance Sampling

- What is the probability of observing a certain trajectory on random variables $a_k, s_{k+1}, a_{k+1}, \dots, s_T$ starting in s_k while following π ?


$$P[a_k, s_{k+1}, a_{k+1}, \dots, s_T | s_k, \pi] = \prod_k^{T-1} \pi(a_k | s_k) p(s_{k+1} | s_k, a_k)$$

- Above p is the state-transition probability.
- Definition: *Importance sampling ratio*
The relative probability of a trajectory under the target and behavior policy, the importance sampling ratio, from sample step k to T is:

$$\rho_{k:T} = \frac{\prod_k^{T-1} \pi(a_k | s_k) p(s_{k+1} | s_k, a_k)}{\prod_k^{T-1} b(a_k | s_k) p(s_{k+1} | s_k, a_k)} = \frac{\prod_k^{T-1} \pi(a_k | s_k)}{\prod_k^{T-1} b(a_k | s_k)}$$

Importance Sampling for Off-Policy Monte-Carlo

- Use returns generated from b (behavior policy) to evaluate π (optimal policy)
- Weight return G_t according to similarity between policies
- Multiply importance sampling corrections along whole episode

return value if have followed π 

$$G_t^{\pi/b} = \frac{\pi(a_t|s_t)\pi(a_{t+1}|s_{t+1})}{b(a_t|s_t)b(a_{t+1}|s_{t+1})} \cdots \frac{\pi(a_T|s_T)}{b(a_T|s_T)} G_t$$
$$E_{X \sim Q} \left[\frac{P(X)}{Q(X)} f(X) \right]$$

- Update value towards corrected return

$$V(s_t) \leftarrow V(s_t) + \alpha(G_t^{\pi/b} - V(s_t))$$

- Cannot use if b is zero when π is non-zero
- Importance sampling can dramatically increase variance ($X = G_t^{\pi/b}$)


$$\text{Var}[X] = E[X^2] - E[X]^2$$

Q Learning

- We now consider off-policy learning of action-values $Q(s,a)$
- **No** importance sampling is required
- Next action is chosen using behavior policy $a_t \sim b(.|s_t)$
- But we consider **alternative successor action** $a' \sim \pi(.|s_t)$
- And update $Q(s_t, a_t)$ towards value of alternative successor action

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_{t+1} + \gamma Q(s_{t+1}, a') - Q(s_t, a_t))$$

- If $a' \sim b(.|s_t)$, becomes Sarsa



Off-Policy Control with Q-Learning


- We now allow both behavior and target policies to improve
- The behavior policy b is e.g. ϵ -greedy w.r.t. $Q(s,a)$
- The target policy π is greedy w.r.t. $Q(s,a)$

$$\pi(S_{t+1}) = \operatorname{argmax}_{a'} Q(S_{t+1}, a')$$

- The Q-learning target then simplifies

$$\begin{aligned} R_{t+1} + \gamma Q(S_{t+1}, A') &= R_{t+1} + \gamma Q(S_{t+1}, \operatorname{argmax}_{a'} Q(S_{t+1}, a')) \\ &= R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a') \end{aligned}$$

Q-Learning Control Algorithm



- Q-learning:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t))$$

- Recall SARSA

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_{t+1} + \gamma Q(s_{t+1}, a') - Q(s_t, a_t))$$


- Q-learning control converges to the optimal action-value function, $Q(s, a) \rightarrow Q_*(s, a)$

Q-Learning for Off-Policy Control


-
- 1: Initialize $Q(s, a), \forall s \in S, a \in A$ $t = 0$, initial state $s_t = s_0$
 - 2: Set π_b to be ϵ -greedy w.r.t. Q
 - 3: **loop**
 - 4: Take $a_t \sim \pi_b(s_t)$ // Sample action from policy
 - 5: Observe (r_t, s_{t+1})
 - 6: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$
 - 7: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 8: $t = t + 1$
 - 9: **end loop**
-

ϵ -greedy policy improvement

greedy policy



Q-Learning vs Sarsa




R. Sutton and A. Barto, Reinforcement Learning: An Introduction, 2018

- But the reason that SARSA took safest path is because the policy of SARSA is the epsilon greedy where *epsilon percent* of the time the agent took random walk.
- This means it is not safe at all to walk close to the cliff
- To avoid the big punishment of agent falling off the cliff, accept the small punishment of long traveling instead.

Double Q-Learning

- Q-Learning performs poorly in some environments because of overestimation of Q values
- Overestimation caused by $\max_{a'} Q(s_{t+1}, a')$ in the following

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(R + \gamma \max_{a'} Q(s_t, a') - Q(s_t, a_t))$$



No	c ₁	c ₂	c ₃	Avg(c _i)	Max(c _i)
1	-0.4	0.2	-0.2	-0.13	0.2
2	0.1	-0.3	0	-0.06	0.1

$$q(C, b) \leftarrow q(C, b) + \alpha(r + \gamma \max_{c_i} q(D, c_i) - q(C, b))$$

$$q(C, b) \leftarrow q(C, b) + (0 + \max\{-0.4, 0.2, -0.2\} - q(C, b)) = q(C, b) + (0.2 - q(C, b))$$

Double Q-Learning

- The proposed solution is to maintain two Q-value functions q_A and q_B , each one gets update from the other for the next state.
- Use q_A to find best action a_* (maximum q)

$$a_* = \underset{a}{\operatorname{argmax}} q_A(s_{t+1}, a)$$

- Use the action value of a_* from q_B

- **Double Q Learning**

$$q_A(s_t, a_t) \leftarrow q_A(s_t, a_t) + \alpha(r_{t+1} + \gamma q_B(s_{t+1}, a_*) - q_A(s_t, a_t)) \quad \text{where } a_* = \underset{a}{\operatorname{argmax}} q_A(s_{t+1}, a)$$

Double Q-Learning

Algorithm 1 Double Q-learning

```
1: Initialize  $Q^A, Q^B, s$ 
2: repeat
3:   Choose  $a$ , based on  $Q^A(s, \cdot)$  and  $Q^B(s, \cdot)$ , observe  $r, s'$ 
4:   Choose (e.g. random) either UPDATE(A) or UPDATE(B)
5:   if UPDATE(A) then
6:     Define  $a^* = \arg \max_a Q^A(s', a)$ 
7:      $Q^A(s, a) \leftarrow Q^A(s, a) + \alpha(s, a) (r + \gamma Q^B(s', a^*) - Q^A(s, a))$ 
8:   else if UPDATE(B) then
9:     Define  $b^* = \arg \max_a Q^B(s', a)$ 
10:     $Q^B(s, a) \leftarrow Q^B(s, a) + \alpha(s, a) (r + \gamma Q^A(s', b^*) - Q^B(s, a))$ 
11:   end if
12:    $s \leftarrow s'$ 
13: until end
```
