



Deep Neural Network & Gradient Methods

Neural Network Representation

- An ANN is composed of processing elements, organized in different ways to form the network's structure.
- Each element receives inputs, processes inputs and delivers a single output.
- The input can be raw input data or the output of other elements. The output can be the final result (e.g. 1 means yes, 0 means no) or inputs to other elements.



Gradient Descent Optimization Methods

- Introduction of topics in gradient descent when training Deep Learning
- Contents
 - Loss Functions
 - Activation Functions
 - Variants of Gradient Method
 - Advanced Methods in Gradient Update
 - Weight Initialization
 - Other Tips

Gradient Descent Method

- Deep learning methods are based on gradient method
- Therefore, many work has been done to improve gradient method
 - 1) Improve learning rate: AdaGrad, RMSProp, AdaDelta, Adam, etc
 - 2) Improve error function: MSE, cross-entropy, etc
 - 3) Improve activation function: ReLU, Leaky ReLU, etc
 - 4) Improve by adding additional terms: Regularization, Momentum, NAG, etc
 - 5) etc

The diagram illustrates the gradient descent update rule: $w \leftarrow w - \eta \frac{\partial}{\partial w} f(w) \pm \alpha$. The components are annotated as follows:

- 1)** Points to the learning rate η .
- 2), 3)** Points to the function $f(w)$.
- 4)** Points to the additional term $\pm \alpha$.

The terms η , $f(w)$, and $\pm \alpha$ are circled in green.

Loss Functions

1. Mean Squared Error (MSE)

$$\frac{1}{n} \sum_i (p_i - y_i)^2$$


- Very popular. Sometime use log form
- Usually better than MAE
- Saturates when using with sigmoid activation function
- Both for **classification** and **regression**
- Without $1/n$, it becomes L2 regularizer

Loss Functions

2. Cross entropy(CE)


$$H(p, y) = \sum_i -(p_i \log y_i)$$

- Binary CE & multi-class CE
- Average of information of p w.r.t. y
- It penalizes heavily for being *very confident* and *very wrong*
- Default loss function to use for classification problems
- Good for classification of small number of class values
- Good regardless of activation function (most popular)
- Faster than MSE



Activation Functions

1. Sigmoid




Takes a real-valued number and “squashes” it into range between 0 and 1.

$$\mathbb{R}^n \rightarrow [0,1]$$

- When activation value is near 0 or 1, the gradient is almost zero, causing **vanishing gradient problem**
- If the initial weights are too large then most neurons would become saturated and the network will barely learn.
- Slow in convergence
- If the data coming into a neuron is always positive, then the gradient on the weights, during backpropagation, will become either all be positive, or all negative (Not zero-centered)
- Computationally expensive
- Can be used in output layer of classification (ranges between 0 and 1)
- It is especially used for models where we have to predict the probability as an output
- The function is monotonic but function's derivative is not.

Activation Functions

2. Tanh




Takes a real-valued number and “squashes” it into range between -1 and 1.

$$\mathbb{R}^n \rightarrow [-1, 1]$$

- Very similar to sigmoid
- Can be used in output layer of classification (ranges between -1 and 1)
- Like sigmoid, tanh neurons can saturate
- Unlike sigmoid, output is zero-centered
- Tanh is a scaled sigmoid: $\tanh(x) = 2\text{sigm}(2x) - 1$
- In practice, the tanh non-linearity is preferred to the sigmoid nonlinearity
- Gradient is stronger than sigmoid which makes the learning faster

Activation Function

3. ReLU (Rectified Linear Unit)



Takes a real-valued number and thresholds it at zero


$$f(x) = \max(0, x)$$

$$R^n \rightarrow R_+^n$$

- It was found to greatly accelerate (e.g. a factor of 6) the convergence of stochastic gradient descent compared to the sigmoid/tanh functions
- Most popular and easy to implement
- Can use in hidden layer, not in output layer
- In output layer, use sigmoid/tanh (binary) or softmax (multi-class) for classification and linear function for regression problem
- Computationally efficient
- Some ReLU units simply die during training. (e.g., large negative weights, η is high)
- Sometimes, ReLU blows up the activation

Activation Functions

4. Softmax



- Special function on **last layer**
- Use in multi-class classification problem
- # of input = # of output
- Squashes a C -dimensional vector \mathbf{O} of arbitrary real values to a C -dimensional vector of real values in the range $(0, 1)$ that add up to 1.
- Each value ranges between 0 and 1 and the sum of all values is 1 so can be used to model probability distributions. Turns the output into a probability distribution on classes.
- Mimics one-hot-encoding
- Only used in the output layer rather than throughout the network
- For multi-label classification, never use softmax. Use sigmoid instead

Variants of Gradient Descent

- Three variants of gradient descent, which differ in how much data we use to compute the gradient of the loss(objective) function.
- Depending on the amount of data, we make a trade-off between the accuracy of the parameter update and the time it takes to perform an update.

1. Batch gradient descent

- Computes the gradient of the cost function w.r.t. to the parameters for the entire training dataset
- Very slow and is intractable for datasets that don't fit in memory.

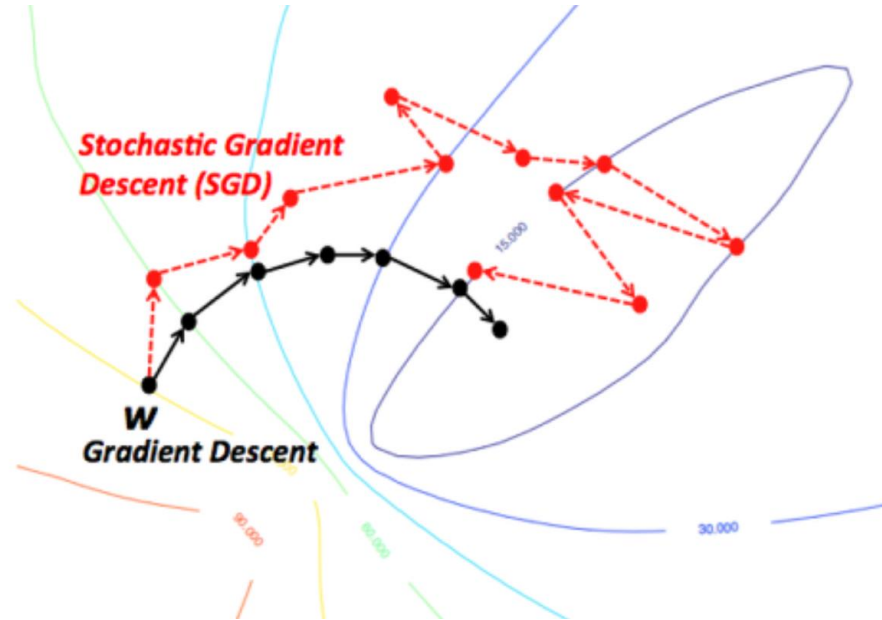
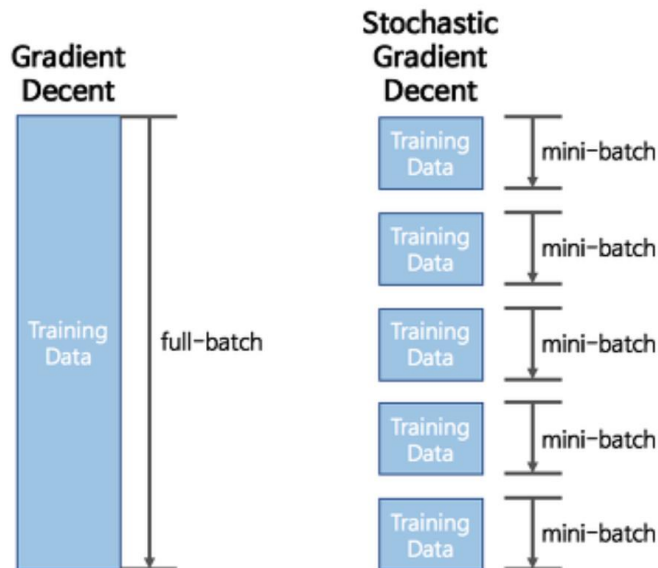
2. Stochastic gradient descent (SGD)

- Performs a parameter update for *each* training example
- Much faster and can also be used to learn online
- Performs frequent updates with a high variance that cause the objective function to fluctuate heavily

Variants of Gradient Descent

3. Mini-batch gradient descent

- Takes the best of both worlds and performs an update for every mini-batch of training examples
- Reduces the variance of the parameter updates, which can lead to more stable convergence
- Can make use of highly optimized matrix optimizations common to state-of-the-art deep learning libraries that make computing the gradient w.r.t. a mini-batch very efficient.
- Common mini-batch sizes range between 50 and 256, but can vary for different applications.



Methods in Learning Rate

$$w_{t+1} = w_t - \eta \nabla_w J(w_t)$$

1. AdaGrad

- Maintains a per-parameter learning rate that improves performance on problems with sparse gradients (e.g. natural language).
- **Low learning rates** for parameters with **frequently occurring features**, and **high learning rates** for parameters with **infrequent features**
- Well-suited for dealing with sparse data
- G_t : a diagonal matrix where each **diagonal element is the sum of the squares of the gradients** w.r.t. w up to time step t . (ϵ : a smoothing term)
- Without the square root operation, the algorithm performs much worse.
- The accumulated sum keeps growing, which in turn causes the learning rate to shrink and eventually become infinitesimally small

$$G_t = G_{t-1} + (\nabla_w J(w_t))^2$$

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \nabla_w J(w_t)$$

Methods in Learning Rate

2. RMSProp

- Adaptive learning rate method proposed by Geoff Hinton in his Lecture
- Maintains per-parameter learning rates that are adapted based on the average of recent magnitudes of the gradients for the weight (e.g. how quickly it is changing).
- The algorithm does well on online and non-stationary problems (e.g. noisy)
- RMSprop as well divides the learning rate by an exponentially decaying average of squared gradients.
- Hinton suggests γ to be set to 0.9, while a good default value for the learning rate η is 0.001

$$G_t = \gamma G_{t-1} + (1 - \gamma) (\nabla_w J(w_t))^2$$

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \nabla_w J(w_t)$$

Methods in Learning Rate

3. AdaDelta

- An extension of Adagrad that seeks to reduce its monotonically decreasing learning rate
- Replace the diagonal matrix G_t with the decaying average over past squared gradients (same as RMSProp)
- In addition, learning rate is replaced by a term of exponentially decaying average squared parameter updates
- No need to define learning rate

$$G_t = \gamma G_{t-1} + (1 - \gamma)(\nabla_w J(w_t))^2 \quad S_t = \gamma S_{t-1} + (1 - \gamma)(\Delta w_t)^2$$

$$\Delta w_t = \frac{\sqrt{S_{t-1} + \epsilon}}{\sqrt{G_t + \epsilon}} \nabla_w J(w_t)$$

$$w_{t+1} = w_t - \Delta w_t$$

Methods in Learning Rate

4. Adam(Adaptive Moment Estimation)

- Moment: n-th moment of a random variable is defined as the expected value of that variable to the power of n. ($m_n = E[X^n]$)
- First moment ($E[X]$) is mean : $E[X]$
- (Since $Var(X) = E[X^2] - E[X]^2$) Second moment ($E[X^2]$) is uncentered variance (meaning we don't subtract the mean during variance calculation).
- Computes the decaying averages of past and past squared gradients m_t and v_t respectively
- m_t and v_t are estimates of the first moment (the mean) and the second moment (the uncentered variance) of the gradients respectively

Methods in Learning Rate

Adam(Adaptive Moment Estimation)

- Combines RMSProp and Momentum

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla_w J(w_t) \quad : \text{Momentum method}$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla_w J(w_t))^2 \quad : \text{RMSProp method}$$

- Use the following unbiased estimate

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

the algorithm calculates an exponential moving average of the gradient and the squared gradient, and the parameters beta1 and beta2 control the decay rates of these moving averages.

moment estimates are biased towards zero. This bias is overcome by first calculating the biased estimates before then calculating bias-corrected estimates.

- The authors propose default values of 0.9 for β_1 , 0.999 for β_2 , and 10^{-8} for ϵ .
- Suggested as the default optimization method for deep learning applications.

Methods in Momentum


Traditional
momentum

$$w_{t+1} = w_t + M_t$$

$$M_t = \alpha M_{t-1} - \eta \nabla_w J(w_t)$$

Nesterov Accelerated Gradient (NAG)


- Have a smarter ball that knows to slow down before the hill slopes up again
- NAG first makes a big jump in the direction of the previous gradient (green vector), measures the gradient (red vector), which results in the final NAG (blue vector).
- This anticipatory update prevents us from going too fast and results in increased responsiveness




$$w_{t+1} = w_t + M_t$$

$$M_t = \alpha M_{t-1} - \eta \nabla_w J(w_t + \alpha M_{t-1})$$


Nesterov Accelerated Gradient (NAG)




Momentum



Momentum



NAG



$$w_{t+1} = w_t + M_t$$

$$M_t = \alpha M_{t-1} - \eta \nabla_w J(w_t + \alpha M_{t-1})$$

Weight Initialization

- Proper initialization of parameters is important
- Never initialize all weight to 0
- Gaussian Initialization with mean=0, sd=1

1) LeCun Initialization

(n_{in} : number of neurons feeding into it, n_{out} : number of neurons the result is fed to)

* LeCun Normal Initialization

$$W \sim N(0, Var(W))$$

$$Var(W) = \sqrt{\frac{1}{n_{in}}}$$

* LeCun Uniform Initialization

$$W \sim U\left(-\sqrt{\frac{1}{n_{in}}}, +\sqrt{\frac{1}{n_{in}}}\right)$$

Weight Initialization

2) Xavier Initialization

* Xavier Normal Initialization

$$W \sim N(0, Var(W))$$
$$Var(W) = \sqrt{\frac{2}{n_{in} + n_{out}}}$$

* Xavier Uniform Initialization

$$W \sim U(-\sqrt{\frac{6}{n_{in} + n_{out}}}, +\sqrt{\frac{6}{n_{in} + n_{out}}})$$

3) He Initialization

* He Normal initialization

$$W \sim N(0, Var(W))$$
$$Var(W) = \sqrt{\frac{2}{n_{in}}}$$

* He Uniform initialization

$$W \sim U(-\sqrt{\frac{6}{n_{in}}}, +\sqrt{\frac{6}{n_{in}}})$$

- Xavier is effective when using sigmoid & tanh and He is effective in ReLU

Batch Normalization

- In general, Gradient descent converges much faster with feature scaling than without it.

$$h_1 = \sigma(w_1 X), h_2 = \sigma(w_2 h_1) = \sigma(w_2 \sigma(w_1 X)), h_3 = \dots$$

- Internal covariate shift
- Batch Normalization (BN) is a normalization method/layer for neural networks
- It consists of **normalizing activation vectors from hidden layers** using the first and the second statistical moments (mean and variance) of the current batch.
- This normalization step is applied right before (or right after) the nonlinear function.

Batch Normalization

- Whitening method

$$\mu_B \leftarrow \frac{1}{m} \sum_{i=1}^m x_i, \quad \sigma_B^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

1) Using whitening method, *bias* parameter is gone.


$$x' = \text{weight} * x + \text{bias}$$

2) loss of non-linearity



- Batch Normalization adds another layer of computation

$$y_i \leftarrow \gamma \hat{x}_i + \beta$$



Regularization

Regularization

- The most widely used method against overfitting in machine learning
- **Regularization**: Machine learning technique that constrains our optimization problem to discourage complex models
- Occam's razor
- Improve generalization of our model on unseen data
- Reduce the complexity of model
 - VC dimension (why not use it?), Number of parameters,
- Very important topic in machine learning

Regularization

- Idea: A new regularization term (regularizer) is added to loss (objective) function
- Loss function with regularization term is modified as follows
 - $J(w)$: loss function, λ : regularization constant
 - $R(w)$: **regularization term** (aka **regularizer**). Can be regarded as the **complexity** of model

$$J_{reg}(w) = J(w) + \lambda * R(w)$$

- Gradient descent method minimizes $J_{reg}(w)$ instead of $J(w)$
- As the value of λ rises, it reduces the value of w 's and thus reducing the variance
- Goal of regularization term ($R(w)$) is to make model parameters simpler, specifically
 - reduces the number of parameters (**L1/Lasso regularization**) or
 - makes the values of parameters smaller (**L2/Ridge regularization**)

Regularizer

- A regularizer is an additional criteria to the loss function to make sure that we don't overfit
- It's called a regularizer since it tries to keep the parameters more normal/regular.
 - significantly reduces the variance of the model, without substantial increase in its bias
- Generally, we don't want huge weights
- If weights are large, a small change in a feature can result in a large change in the prediction. Also gives too much weight to any one feature
- Might also prefer weights of 0 for features that aren't useful
- How do we encourage small weights? or penalize large weights?

Common Regularizers

- From $J_{reg}(w) = J(w) + \lambda * R(w)$

L1 (Lasso) Regularization

- Regularizer, $R(w)$, is given as the sum of weights

$$R(w) = \sum_i |w_i|$$

$$J_{reg}(w) = J(w) + \lambda \sum_i |w_i|$$

- Forces less important parameters be zero
- Reduces the number of features automatically (feature selection)
- Regularization constant (λ , weight decay value) determines how dominant regularization is during gradient computation
- Regularizer is not differentiable
- Generate multiple solutions
- Also known as **Lasso** regularization

Common Regularizers

L2 (Ridge) Regularization

- L2 (Ridge) regularization is given as the sum of weight squares

$$R(w) = \sum_i w_i^2$$

$$J_{reg}(w) = J(w) + \lambda \sum_i w_i^2$$


- Regularization term penalizes big weights
- Big weight decay coefficient \rightarrow big penalty to big weights
- Squared weights penalizes large values more
- Generate a unique solution

Overfitting

Overfitting

- Every machine learning has the issue of overfitting problem
- Occam's razor: prefer simpler model.
- In decision tree, we use pruning technique to avoid overfitting problem
 - Makes the decision tree simpler
- Overfitting is a serious problem in deep learning since it has many layers with too many parameters
- Methods used in Deep Learning (or neural network) to avoid overfitting
 - 1) Regularization
 - 2) Dropout
 - 3) Data Augmentation
 - 4) Early Stop
 - 5) etc

Overfitting




Underfitting

Model does not have capacity to fully learn the data

Ideal fit

Overfitting


Too complex, extra parameters, does not generalize well



Overfitting: Learned hypothesis may **fit** the training data very well, even outliers (**noise**) but fail to **generalize** to new examples (test data)


Dropout

- During training, randomly set some activations to 0 during training
- Each unit retained with fixed probability p , independent of other units
- **Hyper-parameter** p to be chosen (tuned)
- Typically drop 50% of activations in layer
- Forces network to not rely on any 1 node
- Select different dropout nodes at each mini-batch
- May use different p value in layerwise



Early Stopping

- Use validation error to decide when to stop training
- Stop training before we have a chance to overfit
- Stop when monitored quantity has not improved after n subsequent epochs
- n is called patience
- generally this is not recommended by some people



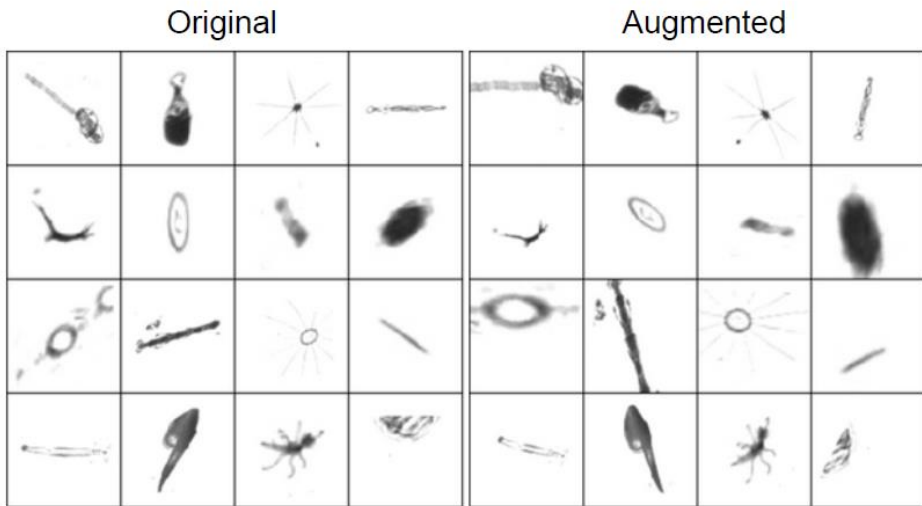
Data Augmentation

- In practice, the amount of data we have is limited.
- One way to get around this problem is to create fake data and add it to the training set.
- Dataset augmentation has been a particularly effective technique for a specific classification problem: object recognition.
- Images are high dimensional and include an enormous variety of factors of variation, many of which can be easily simulated.
- Operations like translating the training images a few pixels in each direction can often greatly improve generalization

Data Augmentation

- Techniques for Data Augmentation
 - Flip: flip images horizontally and vertically
 - Rotation: image dimensions may not be preserved after rotation
 - Scale: image can be scaled outward or inward
 - Crop: randomly sample a section from the original image. Then resize this section to the original image size
 - Translation: move the image along the X or Y direction (or both).
- (*) Advanced Augmentation Techniques
 - E.g.: Landscapes in different seasons.
 - Use GAN (generative Adversarial Network) to generate new images for data augmentation


Data Augmentation



Pre-processed images (left) and augmented versions of the same images (right).

Source: Ian Goodfellow et al. Deep Learning, MIT press, 2016

Geometry based



Color based



Noise / occlusion



Weather





<https://blog.insightdatascience.com/automi-for-data-augmentation-e87cf692c366?gi=bc0b8b5a353a>

CNN

Drawbacks of Neural Networks

- Little or no invariance to shifting, scaling, and other forms of distortion



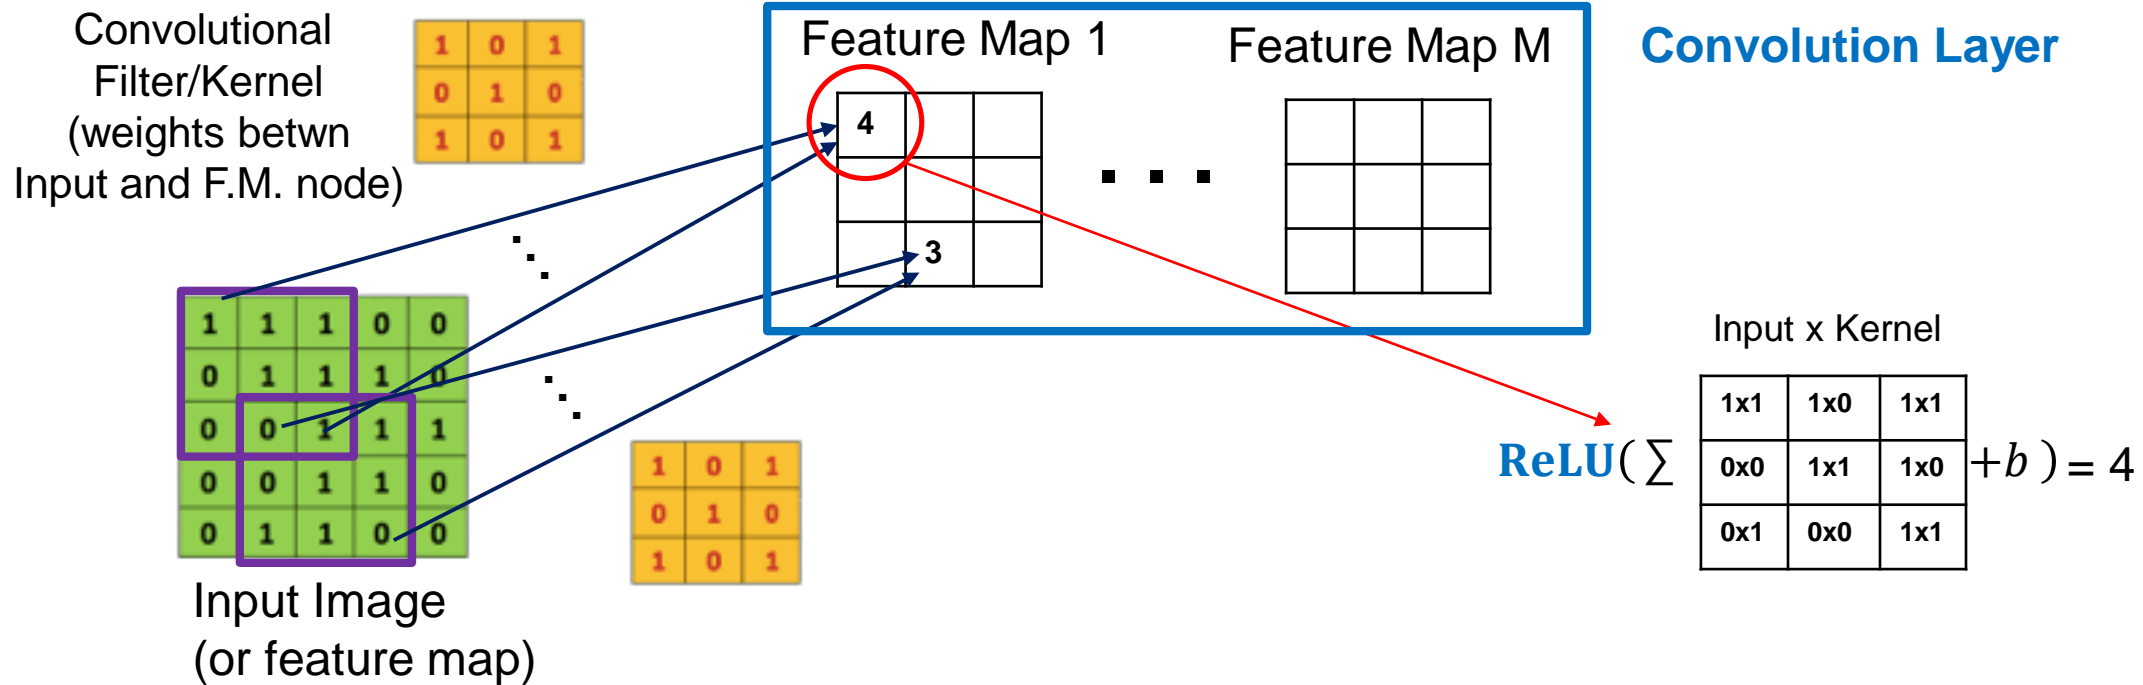
154 input change from 2 shift left
77 : black to white
77 : white to black



Basic Ideas of Convolutional Neural Networks

- In traditional neural networks, **every neuron** in the network is connected to **every neuron** in adjacent layers (fully-connected)
- As we increase the number of hidden layers (deep neural network), the number of parameters (weights) increases exponentially, causing overfitting
- How do we significantly reduce the number of parameters with hidden layers
- Convolutional neural networks(CNN) use a special architecture which is particularly **well-adapted to classify images**.
 - 1) Every node in hidden layer has the **same weight vector** with lower level layer nodes.
 - 2) Every node in hidden layer is connected with a **small portion** of lower level layer
 - 3) The number of nodes in hidden layer is reduced even further by using **pooling layer**

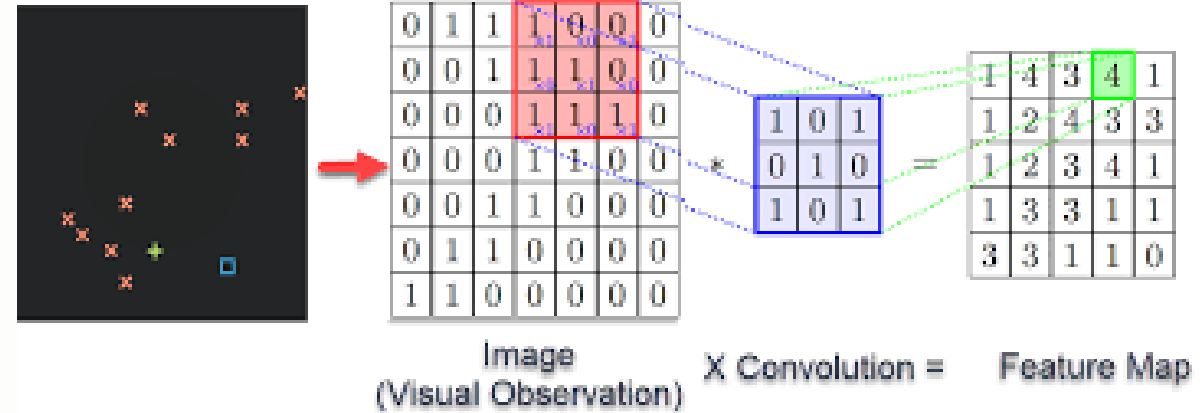

Convolutional Layer



- Convolution filter represents the weight values between Input and feature map
- Every node in the same feature map shares the same convolution filter (**weight sharing**)
- Every node in feature map looks at a small portion of input (**sparse connectivity**)
 - Each node in feature map looks at a different region of input
- Convolution Filters are NOT given by human. It is learned from algorithm

Convolutional Layer


- Convolutional Layer consists of multiple feature maps
- Each feature map has different Convolution Filters
- Each feature map is to detect a certain part of an image (e.g. nose)



<https://www.superdatascience.com/blogs/convolutional-neural-networks-cnn-step-1-convolution-operation>

<https://www.oreilly.com/library/view/learn-unity-ml-agents/9781789138139/16671cc8-0aff-433a-878c-7430be8b9aa1.xhtml>


Convolution Layer (multiple channels)



Sparse Connectivity & Weight Sharing

Hidden or
input layer


Hidden or
output layer



Fully Connected

Input layer or
feature map

Feature map



Sparse Connectivity/
Weight sharing

Role of Feature Maps

- All the neurons in the first hidden layer **detect exactly the same feature**. Informally, think of the feature detected by a hidden neuron as the kind of input pattern: it might be an edge in the image, for instance, or maybe some other type of shape, just at different locations in the input image.
- To do image recognition we'll need **more than one feature map**. And so a complete convolutional layer consists of several different feature maps
- Convolutional networks are well adapted to the translation **invariance of images**
- A big advantage of sharing weights and biases is that it **greatly reduces the number of parameters** involved in a convolutional network


Pooling Layer

- Pooling layers are usually used immediately after convolutional layers. What the pooling layers do is **simplify the information in the output** from the convolutional layer
- A pooling layer takes each feature map output from the convolutional layer and prepares a **condensed feature map**
- A big benefit is that there are many fewer pooled features, and so this helps **reduce the number of parameters** needed in later layers.

Pooling Layer

- For instance, each unit in the pooling layer may summarize 2X2 neurons in the previous layer.
- One common procedure for pooling is known as *max-pooling*. In max-pooling, a pooling unit simply outputs the maximum activation in the 2X2 region
- We can think of max-pooling as a way for the network to ask whether a given feature is found anywhere in a region of the image. It then throws away the exact positional information.


Feature Map




Pooling Layer

Pooling Layer


- The intuition is that once a feature has been found, *its exact location isn't as important* as its rough location relative to other features.
- Max-pooling isn't the only technique used for pooling. Other common approach are known as *Sum pooling & Average Pooling*. Here, instead of taking the maximum activation, we take the sum (or average) of the activations in the region.



Max

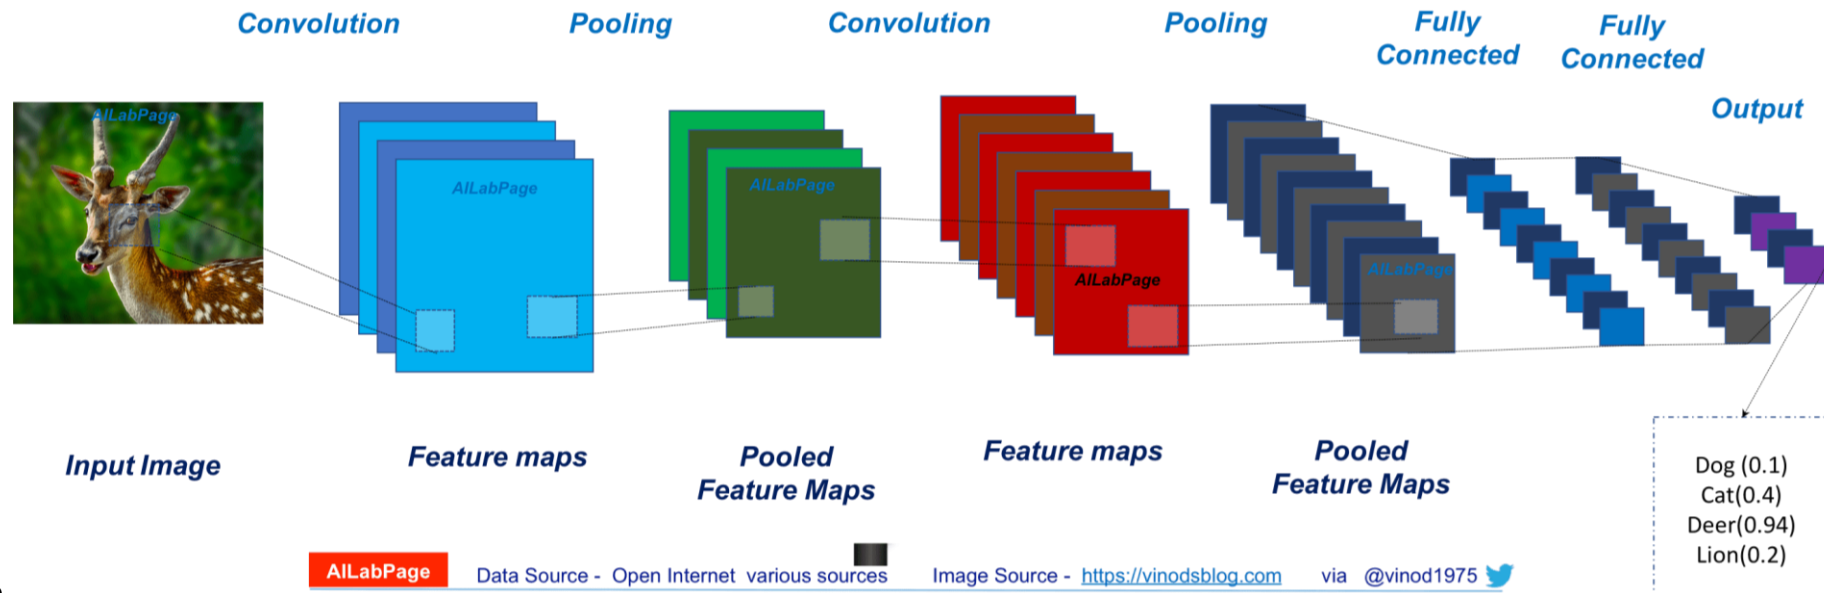


Sum




Putting It All Together

- The **final layer** of connections in the network is a **fully-connected** layer. That is, this layer connects *every* neuron from the max-pooled layer to every one of the output neurons.
- Train the network using **stochastic gradient descent and backpropagation**.
- Need to make few modifications to the backpropagation procedure due to *weight sharing, sparse connectivity, etc*



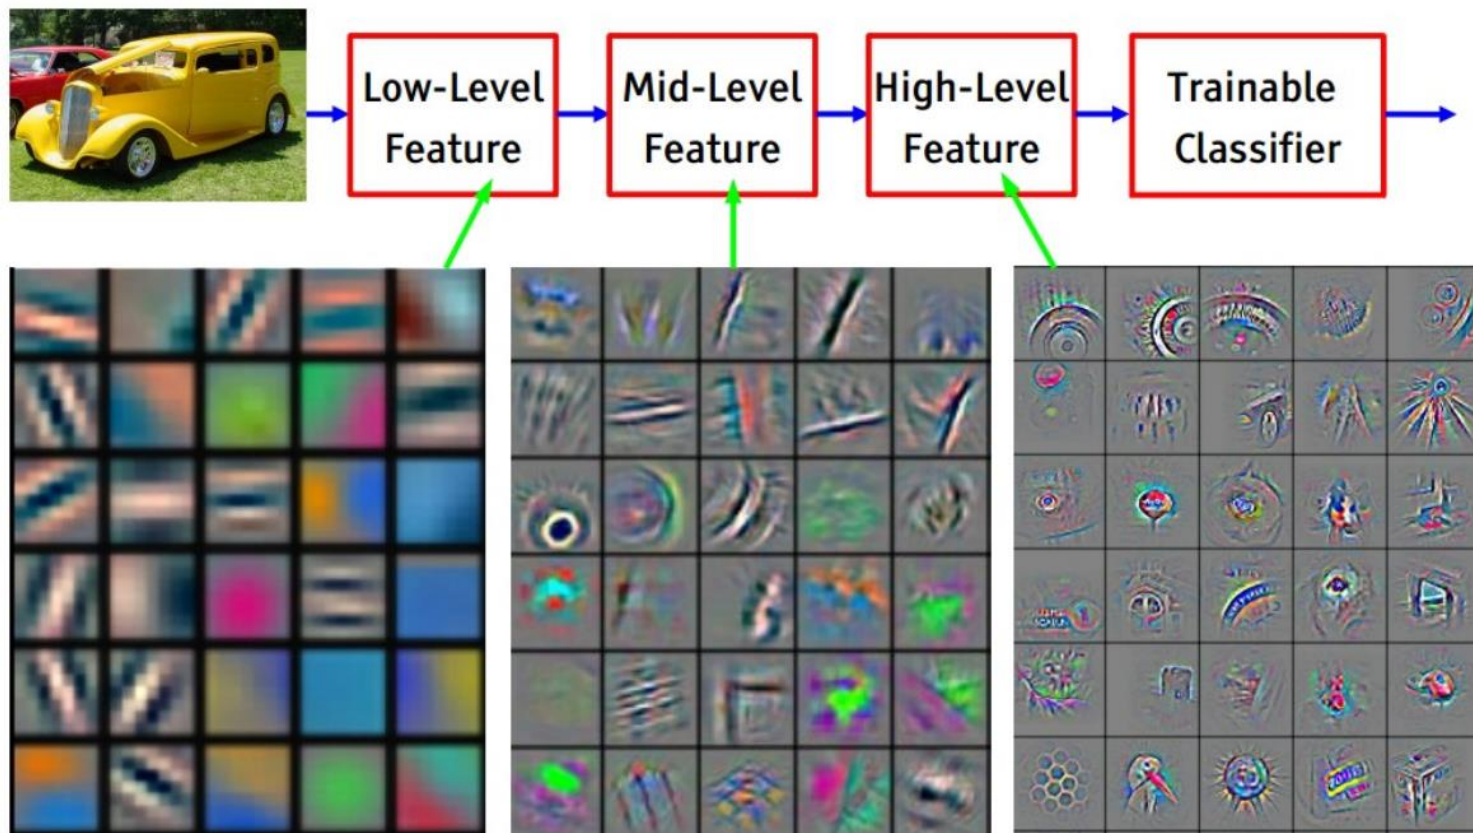
CNN's Topology



<https://www.easy-tensorflow.com/tf-tutorials/convolutional-neural-nets-cnns>

- Convolution and/or pooling can be repeated many times
- Full connection
 - Regular neural network connection
 - No weight sharing

Hierarchical Feature Extraction



RNN

Why do we need RNNs?

- The data we have seen so far is assumed to be i.i.d. (independent and identical distributions)
- Many real data is in temporal order.
 - Frames from video, Words in sentence, Snippets of audio,
- Characteristics of sequential data
 - 1) Current data depends on past data.
 - E.g.: Current word depends on previous words
 - Current frame depends on previous frame
 - 2) Temporal data is variable length
- The limitations of the Neural network (including CNNs)
 - Rely on the assumption of independence among the (training and test) examples.
 - After each data point is processed, the entire state of the network is lost
 - Rely on examples being vectors of fixed length
- We need to model the data with temporal/sequential structures and varying length of inputs and outputs


Recurrent Neural Network

- RNN: Neural network where the current value of node is used in next time step
- Can produce an output at each time step
- Can naturally process variable length temporal data
- Train using back-propagation through time(BPTT)
- Use same set of weights at all time steps

$$\text{NN: } h_t = \sigma(W^{(hx)}x_t)$$


$$\text{RNN: } h_t = \sigma(W^{(hh)}h_{t-1} + W^{(hx)}x_t)$$

Previous state



<https://pbs.twimg.com/media/C2j-8j5UsAACgEK.jpg>

\triangleq

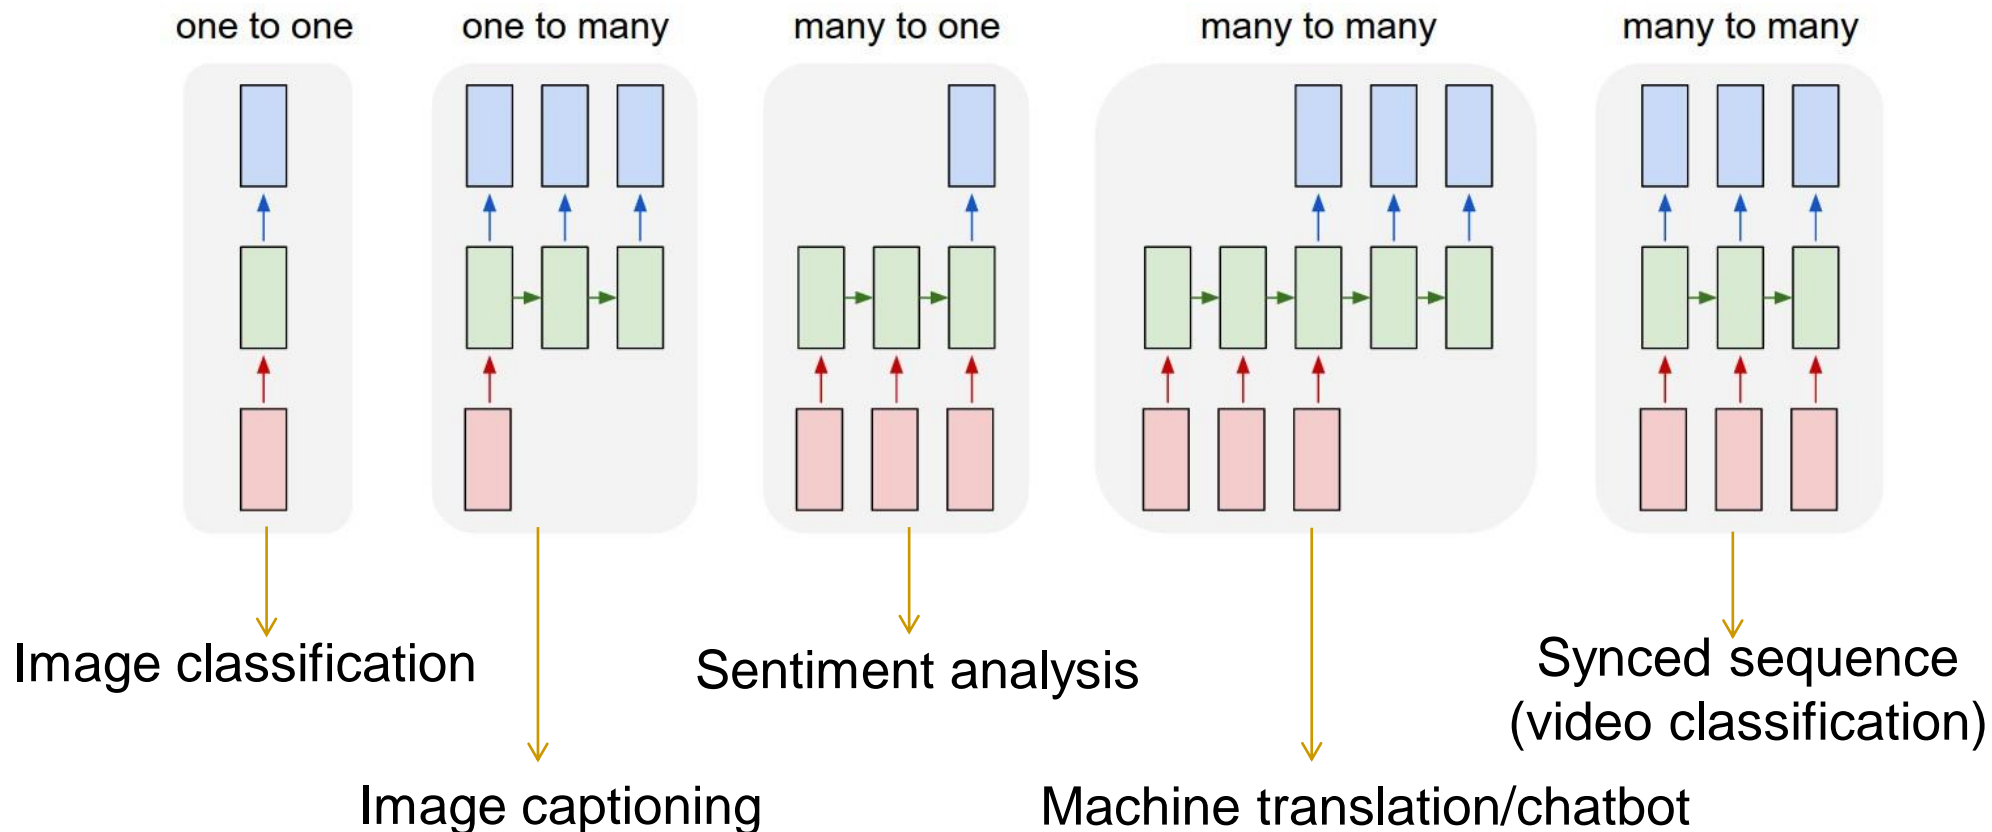


Sharing Weights in RNN

- In CNN, we share the weights across different regions.
- Sharing weights could reduce the number of parameters & improve the performance
- Let's share the weights AGAIN in RNN
 - Share weights across **regions** -> CNN
 - Share weights across **time step** -> RNN
 - Share weights across **tree structure** -> **Recursive NN** (not covered here)
- Sparse connectivity (in CNN) is not used in RNN. RNN is fully connected network.

Sequence to Sequence(Seq2Seq) Learning

- Since RNN can produce an output at each time step, it enables Seq2Seq learning
- **Seq2Seq learning**: given input, algorithm generates a **sequences of output**



Seq2Sequence Learning Example


- e.g.: train RNN to generate “hello” in sequence

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

**Character-level
language model
example**

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

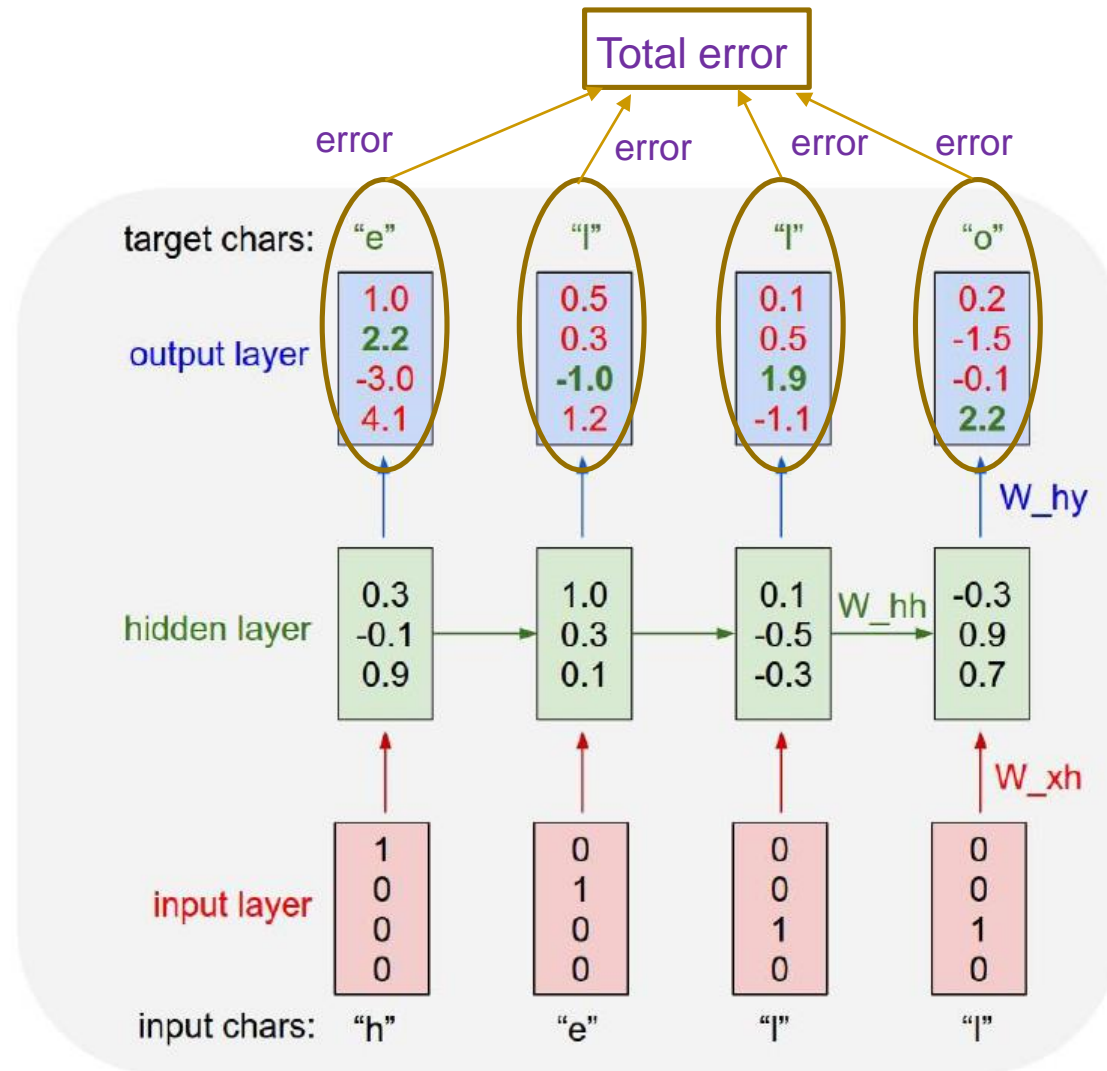


Seq2Sequence Learning Example

Character-level language model example

Vocabulary:
[h,e,l,o]

Example training sequence:
"hello"



Training time:

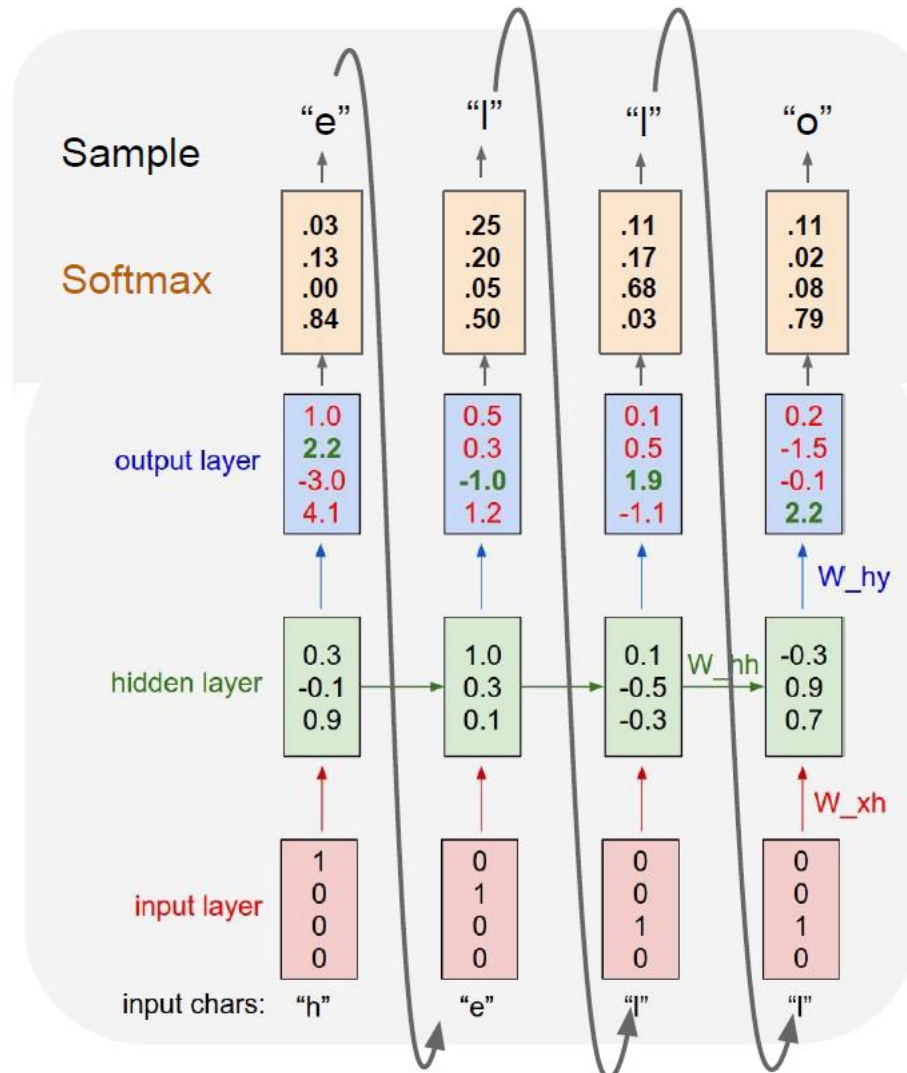
- Use BPTT to train entire RNN network
- (W_{hy} , W_{hh} , W_{xh}) is same in each time step

Seq2Sequence Learning Example

Character-level language model example

Vocabulary:
[h,e,l,o]

Example training sequence:
“hello”

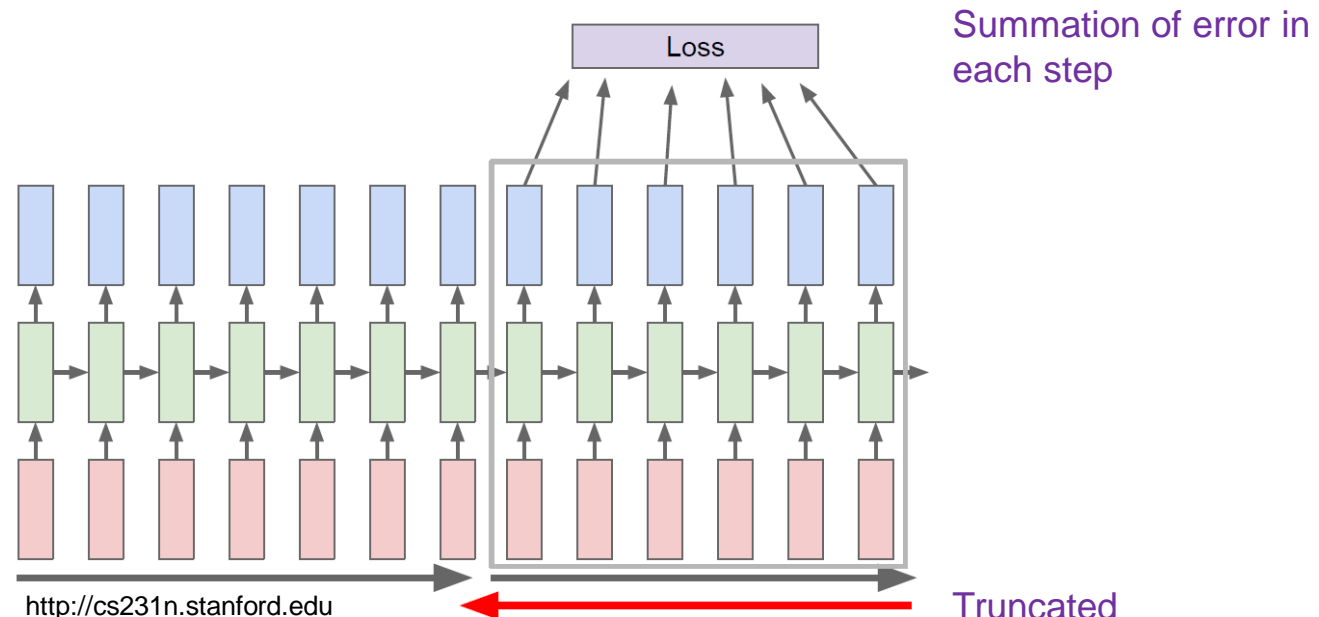


Test time:


- Generate output at each time step
 - pick a seed character sequence
 - generate the next character
 - feed back into the model
 - then the next ...
- This process is basic concept of generative learning
- When generating next sample, use probabilistic sampling from softmax (Not deterministic sampling)
- We assumed to pick 'e' at the first step even though 'o' has the highest prob.

(Truncated) Backpropagation Through Time

- Back Propagation Through Time (BPTT) is used to learn the RNN
- BPTT is an extension of the back-propagation (BP)
- Vanishing Gradients in RNN are Different from the Case in NNs
- If it was just a case of vanishing gradients in NNs, we could just rescale the per-layer learning rate, but that does not really fix the training difficulties
- Can't do that with RNNs because the weights are shared, & total true gradient = sum over different "depths"
- We can't look back forever, and use truncated BPTT



(Truncated) Backpropagation Through Time



- Error function: $J_n = -\sum_k L(y_k, \widehat{y}_k)$
- BPTT also uses gradient descent


\widehat{y}_k : true value
 y_k : prediction
 L : Loss function

- Update $w = w - \eta \frac{\partial J_n}{\partial w}$

Vanishing Gradient in RNN

$$\frac{\partial J_n}{\partial W} = \sum_{k=0}^n \frac{\partial J_n}{\partial y_n} \frac{\partial y_n}{\partial s_n} \frac{\partial s_n}{\partial s_k} \frac{\partial s_k}{\partial W}$$


$\frac{\partial s_n}{\partial s_{n-1}} \frac{\partial s_{n-1}}{\partial s_{n-2}} \cdots \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial s_1} \frac{\partial s_1}{\partial s_0}$



Gradient contributions from “far away” steps become zero, and the state at those steps doesn’t contribute to what you are learning: You end up not learning long-range dependencies.

- Vanishing gradient in RNN means it forgets the past
- It doesn’t remember what has happened in the past

Long Short-Term Memory(LSTM)



- 1) $f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + b_f)$ (forget gate)
- 2) $i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + b_i)$ (input gate)
- 3) $o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + b_o)$ (output gate)
- 4) $g_t = \tanh(W_{xg}x_t + W_{hg}h_{t-1} + b_g)$
- 5) $c_t = c_{t-1} \otimes f_t + g_t \otimes i_t$
- 6) $h_t = \tanh(c_t) \otimes o_t$

Traditional RNNs are a special case of LSTMs: Set the input gate, the forget gate, the output gate to all ones.


- Use **memory** cell to store information at each time step.
- Use “gates” to control the flow of information through the network.
 - **Forget gate**: limit information passed from one cell to the next
 - **Input gate**: how much of the new information will be let through the memory cell.
 - **Output gate**: how much of the information will be passed to expose to the next time step.

Advantages of LSTM

- Non-decaying error backpropagation.
- For long time lag problems, LSTM can handle noise and continuous values.
- No parameter fine tuning.
- Memory for long time periods
- LSTM solves the vanishing gradient and the long memory limitation problem
- LSTM can learn sequences with more than 1000 time steps.


Gated Recurrent Unit (GRU)

- A very simplified version of the LSTM
 - Only two gates: 'update' gate and 'reset' gate
 - Merges forget and input gate into a single 'update' gate
 - Merges cell and hidden state
 - No non-linearity in output value
- Has fewer parameters than an LSTM and has been shown to outperform LSTM on some tasks



1) $r_t = \sigma(W_{xr}x_t + W_{hr}h_{t-1} + b_r)$	(reset gate)
2) $z_t = \sigma(W_{xz}x_t + W_{hz}h_{t-1} + b_z)$	(update gate)
3) $g_t = \tanh(W_{hg}(r_t \otimes h_{t-1}) + W_{xg}x_t + b_g)$	
4) $h_t = h_{t-1} \otimes z_t + g_t \otimes (1 - z_t)$	

GRU



- Compute a reset gate based on current input vector and hidden stat
- Compute an update gate based on current input vector and hidden stat

$$r_t = S \left(W_r \begin{pmatrix} x_t \\ h_{t-1} \end{pmatrix} + b_f \right)$$

$$z_t = S \left(W_z \begin{pmatrix} x_t \\ h_{t-1} \end{pmatrix} + b_f \right)$$

- New memory

$$h'_t = \tanh W \begin{pmatrix} x_t \\ r_t \otimes h_{t-1} \end{pmatrix}$$

- Final memory

$$h_t = (1 - z_t) \ddot{A} h_{t-1} + z_t \ddot{A} h'_t$$