

Advanced Policy Gradient

Contents

- A3C
- MaxEnt RL
- TRPO
- PPO
- DDPG
- TD3

A3C(Asynchronous Advantage Actor Critic)

- Actor-Critic method is an effective, but online learning method
- Not easy to generate many data
- Not easy to train in batch
- Experience replay is one solution, but sometime memory becomes too big.
- Another method of reducing dependency in data is A3C
- Another advantage of A3C is it can be done in parallel.

A3C

- **Asynchronous** stands for the principal difference of this algorithm from DQN, where a single neural network interacts with a single environment.
- In A3C, we've got a global network with multiple agents having their own set of parameters.
- It creates every agent's situation interacting with its environment and harvesting the different and unique learning experience for overall training.
- That deals partially with sample correlation, a big problem for neural networks, which are optimized under the assumption that input samples are independent of each other
- Asynchronous: the algorithm involves executing a set of environments in parallel to increase the diversity of training data, and with gradient updates performed in a Hogwild! style procedure.
- No experience replay is needed, though one could add it if desired.

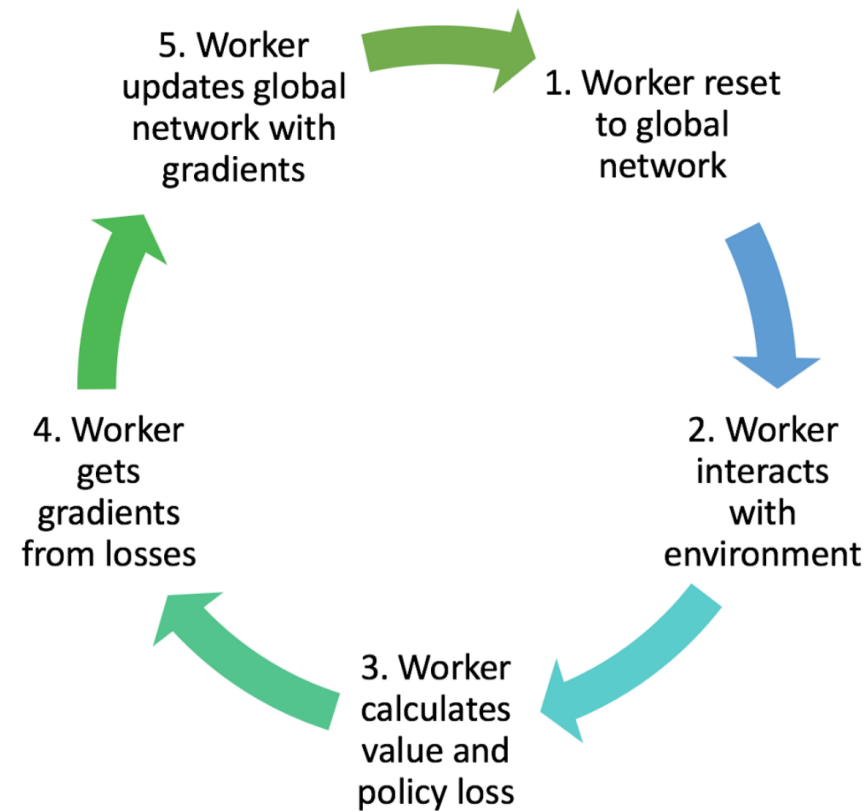

A3C

- A3C idea: Sample for data can be parallelized using several copies of the same agent
 - Use N copies of the agents (workers) working in parallel collecting samples and computing gradients for policy and value function
 - After some time, pass gradients to a global network that updates actor and critic using the gradients of all agents
 - After some time the worker copy the weights of the global network
- This parallelism decorrelates the agents' data, so no Experience Replay Buffer needed
- Even one can explicitly use different exploration policies in each actor-learner to maximize diversity
- Asynchronism can be extended to other update mechanisms (SARSA, Q-learning, etc) but it works better in Advantage Actor Critic setting

A3C



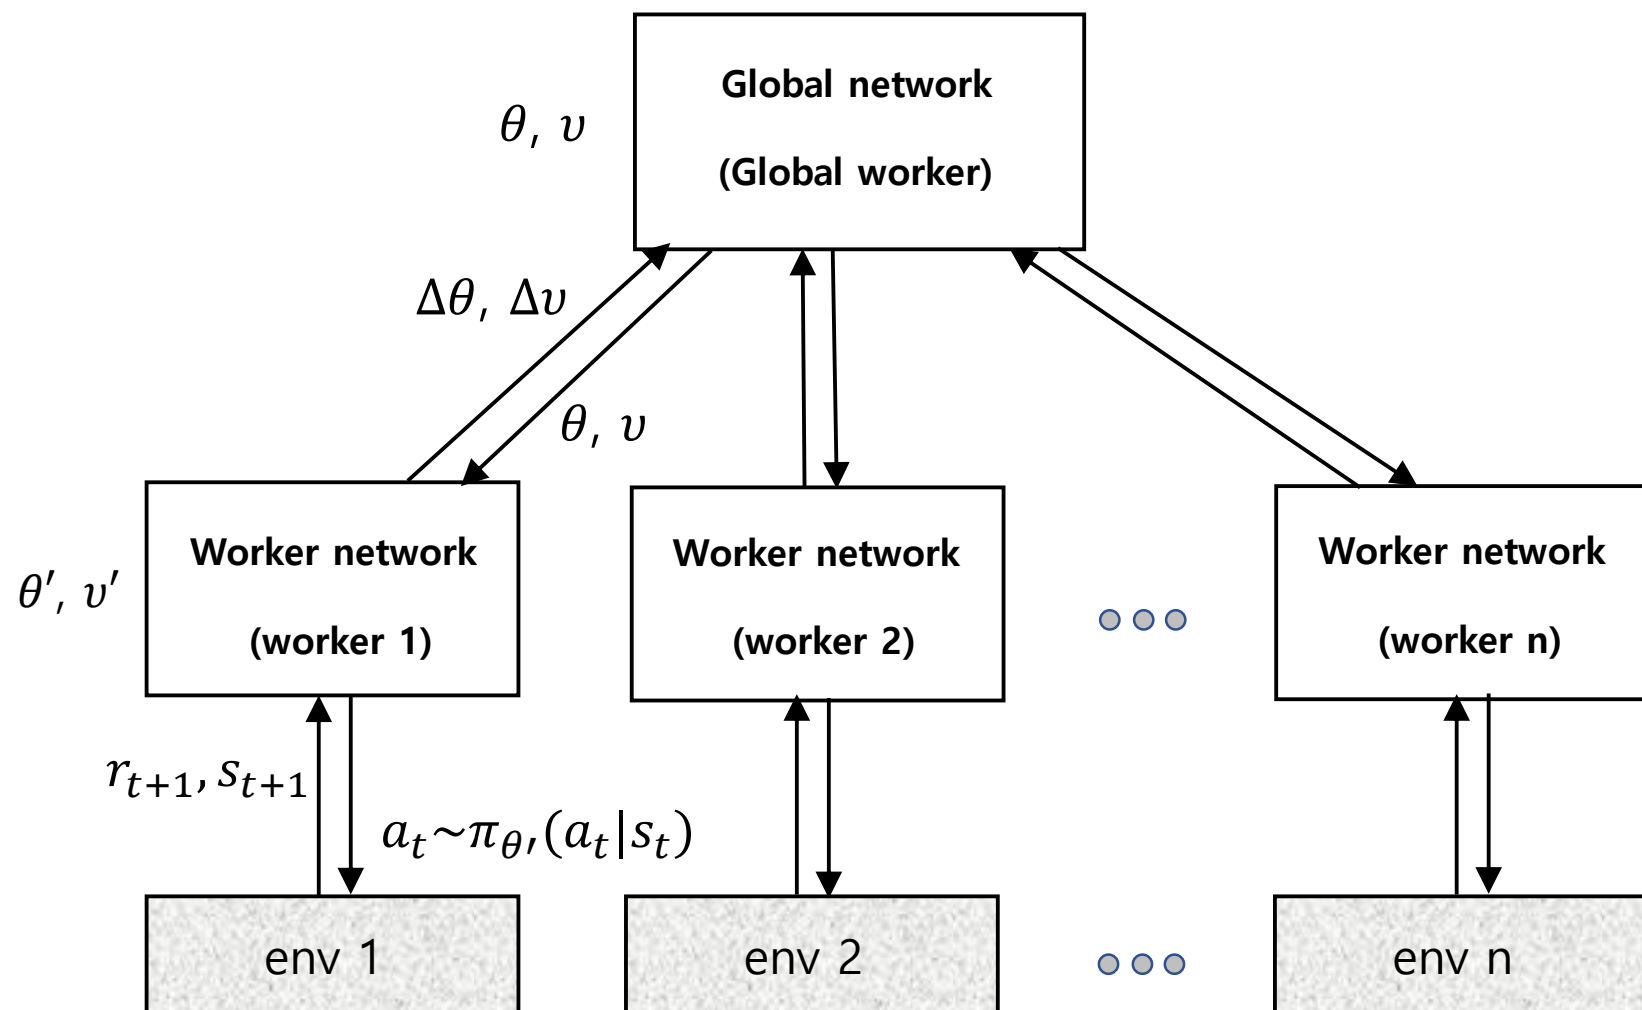
- A3C consists of **multiple independent workers**(agents, networks) with their own weights, who interact with a different copy of the environment in parallel.
- Thus, they can explore a bigger part of the state-action space in much less time.
- The agents (or workers) are trained in parallel and update periodically a global network, which holds shared parameters.
- The updates are not happening simultaneously and that's where the asynchronous comes from.
- After each update, the agents resets their parameters to those of the global network and continue their independent exploration and training for n steps until they update themselves again.

A3C



<https://medium.com/emergent-future/simple-reinforcement-learning-with-tensorflow-part-8-asynchronous-actor-critic-agents-a3c-c88f72a5e9f2>

A3C



Actor Network Update

- Recap: Advantage Actor-Critic (p. 41 in Chapter 9)

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \delta^{\pi_{\theta}}]$$

where $\delta^{\pi_{\theta}} = r + \gamma V_{\pi_{\theta}}(s') - V_{\pi_{\theta}}(s)$

- Therefore, each worker (agent) in A3C computes policy gradient as follows

$$\nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) \delta_t \text{ where } \delta_t = r_{t+1} + \gamma v_{v'}(s_{t+1}) - v_{v'}(s_t)$$

- θ' and v' : parameters of policy network and value network of workers, respectively

- However, A3C uses n-step return (instead of single-step advantage $\delta^{\pi_{\theta}}$)

- n-step return in state s_t : $G_t^{(n)} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n v_{v'}(s_{t+n})$
 - n-step return has the following property

$$G_t^{(n)} = r_{t+1} + \gamma G_{t+1}^{(n-1)}$$

- Using $G_t^{(n)}$, A3C worker agent computes policy gradient as follows

$$\nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) \delta_t^{(n)} \text{ where } \delta_t^{(n)} = G_t^{(n)} - v_{v'}(s_t)$$

Actor Network Update


- A3C uses a trick to compute policy gradient efficiently
- Worker moves from state s_t to s_{t+n} with the following n actions (A3C uses n -step return)
 $(a_t, a_{t+1}, \dots, a_{t+n-1})$

- In each of intermediate state s_{t+k} ($0 \leq k \leq n - 1$), it computes $G_{t+k}^{(n-k)}$
 - $G_{t+k}^{(n-k)}$: value of $(n - k)$ -step return at (intermediate) state s_{t+k}
- For each intermediate state, agent computes policy gradient
 - $\Delta\theta$ is the summation of these gradients, defined as follows

$$\Delta\theta = \sum_{k=0}^{n-1} \nabla_{\theta} \log \pi_{\theta}(a_{t+k}|s_{t+k}) \left(G_{t+k}^{(n-k)} - v_{v'}(s_{t+k}) \right)$$

- Above $\Delta\theta$ is forwarded to global worker to update global network
- After that, worker receives new parameters from global network and starts over
- Each worker works on a little different parameters from global networks.
- As long as learning constant is small, it works fine (*HogWild update*)

Computing Policy Gradients in Intermediate States



$$\Delta\theta = \sum \left\{ \nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) \left(G_t^{(10)} - v_{v'}(s_t) \right) \quad \dots \quad \nabla_{\theta'} \log \pi_{\theta'}(a_{t+9} | s_{t+9}) \left(G_{t+9}^{(1)} - v_{v'}(s_{t+9}) \right) \right\}$$

$$\begin{aligned} \Delta\theta &= \sum_{k=0}^9 \nabla_{\theta'} \log \pi_{\theta'}(a_{t+k} | s_{t+k}) \left(G_{t+k}^{(10-k)} - v_{v'}(s_{t+k}) \right) = \\ &\nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) (r_{t+1} + \gamma r_{t+2} + \dots + \gamma^9 r_{t+10} + \gamma^{10} v_{v'}(s_{t+10}) - v_{v'}(s_t)) + \\ &\quad \dots \\ &\nabla_{\theta'} \log \pi_{\theta'}(a_{t+9} | s_{t+9}) (r_{t+10} + \gamma v_{v'}(s_{t+10}) - v_{v'}(s_{t+9})) \end{aligned}$$

Critic Network Update

- Worker should update critic network as well
- Error function of critic network (with single-step TD(0)) is given as follows (using MSE)

$$E_{\pi_{\theta'}} \left[\left(r_{t+1} + \gamma v_{v'}(s_{t+1}) - v_{v'}(s_t) \right)^2 \right]$$

- Previously, worker computed multi-step return in each intermediate step
- We reuse these multi-step returns in critic network learning
- Therefore, in time step $t + k$, error function is defined as

$$\left(G_{t+k}^{(n-k)} - v_{v'}(s_{t+k}) \right)^2$$

- Therefore, the total error function of critic network is

$$J(v') = \sum_{k=0}^{n-1} \left(G_{t+k}^{(n-k)} - v_{v'}(s_{t+k}) \right)^2$$

- Its derivative is

$$\nabla_{v'} J(v') = \sum_{k=0}^{n-1} \nabla_{v'} \left(G_{t+k}^{(n-k)} - v_{v'}(s_{t+k}) \right)^2$$

A3C

// Assume global shared parameter vectors θ and θ_v and global shared counter $T = 0$

// Assume thread-specific parameter vectors θ' and θ'_v

Initialize thread step counter $t \leftarrow 1$

repeat

Reset gradients: $d\theta \leftarrow 0$ and $d\theta_v \leftarrow 0$.

Synchronize thread-specific parameters $\theta' = \theta$ and $\theta'_v = \theta_v$

$t_{start} = t$

Get state s_t

repeat

Perform a_t according to policy $\pi(a_t|s_t; \theta')$

Receive reward r_t and new state s_{t+1}

$t \leftarrow t + 1$

$T \leftarrow T + 1$

until terminal s_t **or** $t - t_{start} == t_{max}$

$R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, \theta'_v) & \text{for non-terminal } s_t // \text{ Bootstrap from last state} \end{cases}$

for $i \in \{t - 1, \dots, t_{start}\}$ **do**

$R \leftarrow r_i + \gamma R$


Accumulate gradients wrt θ' : $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$

Accumulate gradients wrt θ'_v : $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

end for

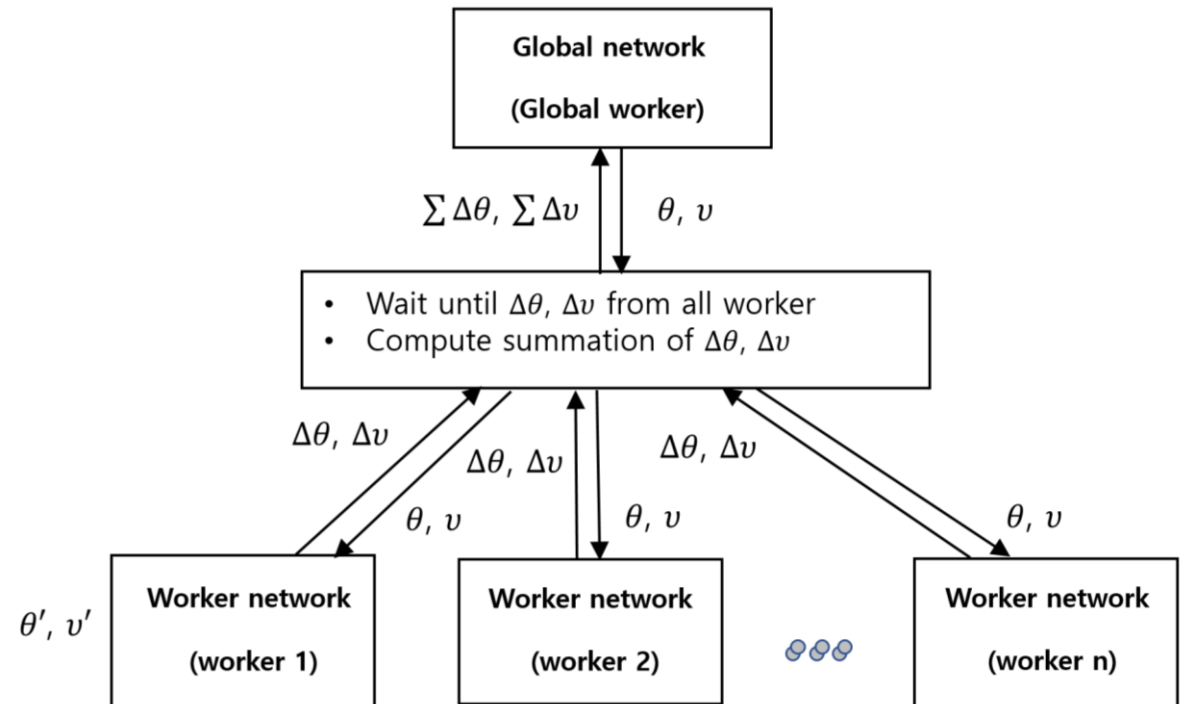
Perform asynchronous update of θ using $d\theta$ and of θ_v using $d\theta_v$.

until $T > T_{max}$


$$G_t^{(n)} = r_{t+1} + \gamma G_{t+1}^{(n-1)}$$

A2C(Advantage Actor Critic)

- In A3C, some agents will be playing with an older version of the parameters.
- Of course, the update may not happen asynchronously but at the same time.
- In that case, we have an improved version of A2C with multiple agents instead of one.
- A2C will wait for all the agents to finish their segment and then update the global network weights and reset all the agents.



MaxEnt RL

- In Maximum Entropy RL, the aim is to learn the optimal policy that can achieve the highest cumulative reward and maximum entropy.
- In Maximum Entropy RL, it enables more exploration and chances to avoid converging to local optima is higher.
- Original objective function

$$E_{\pi_{\theta}} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right]$$

changes to


$$J(\theta) = E_{\pi_{\theta}} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) + \beta H(\pi_{\theta}(\cdot | s_t)) \right]$$

MaxEnt RL

- Why add an entropy regularizer?
- Why is maximizing entropy desirable in RL?

$$H(x) = - \sum_x p(x) \log p(x)$$

$$H(\pi_\theta) = E_{a \sim \pi_\theta} [-\log \pi_\theta(a|s)] = - \sum_a \pi_\theta(a|s) \log \pi_\theta(a|s)$$



Trust Region Policy Optimization(TRPO)

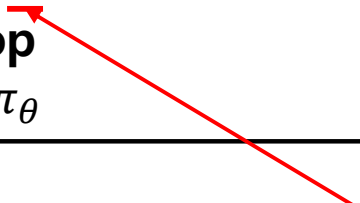
- In supervised learning, even if the parameter moves to a very bad state, it can recover from the learning of the next batch. (Each data is independent of each other by IID)
- However, policy gradient is basically an online learning method.
- Therefore, if you make a mistake in updating the parameters and move to a very bad state (such as setting the wrong step size), learning often becomes almost impossible from the next time.
- There are considerations when defining the confidence interval: 1) whether to define a confidence interval for the change difference in the parameter, 2) or a confidence interval for the change difference in the policy (as determined by the parameter)
- First, defining a confidence interval for a change in parameter values is simple to calculate the confidence interval, but the change in the objective function often changes significantly even with a slight change in the parameter value.
- On the other hand, in general, compared to changes in policy, the objective function changes relatively smoothly.

Motivation – Problem in REINFORCE

```
Initialize policy network  $\pi_\theta$ 
loop
  Generate episodes  $\tau \sim s_0, a_0, r_1, \dots, s_{n-1}, a_{n-1}, r_n$  with  $\pi_\theta$ 
   $\Delta\theta = 0$ 
  for every step  $t = 0, 1, 2, \dots, n - 1$  in episode do
    
$$G_t = \sum_{i=0}^{n-t} \gamma^i r_{t+i}$$

    
$$\Delta\theta \leftarrow \Delta\theta + \gamma^t \nabla_\theta \log \pi_\theta(s_t, a_t) G_t$$

  end for
   $\theta \leftarrow \theta + \alpha \Delta\theta$ 
end loop
return  $\pi_\theta$ 
```



- How to choose the step size(α)?
 - too large? 1) bad policy -> 2) collected data under bad policy
 - too small? cannot leverage data sufficiently
- Can't recover!

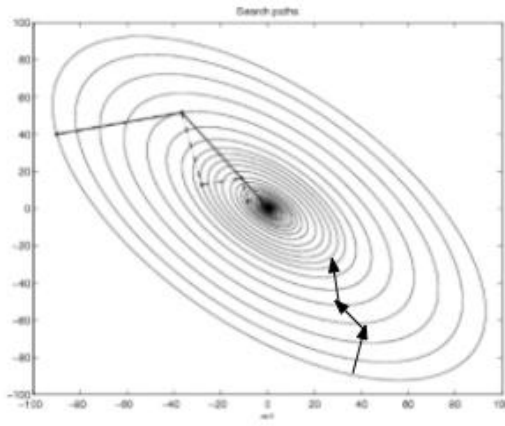
Why Trust Region?




Line search
(like gradient ascent)




Trust region



LINE SEARCH METHOD



TRUST REGION METHOD



$$\theta_k \rightarrow \theta'_{k+1}$$

(Once you fail, you can't recover)

- Instead of following the derivative,
- 1) define a trust(safe) region at current point and
 - 2) move to the optimal point within the trust region

TRPO

- MDP model $\langle S, A, P, R, \rho_0, \gamma \rangle$ where ρ_0 is prob. of starting states
 - $\eta(\pi)$: objective function of policy π
 - θ and $\tilde{\theta}$: current parameter and new parameter of policy network, respectively
 - π, π_θ : current policy
 - $\tilde{\pi}, \pi_{\tilde{\theta}}$: new policy

- Objective function

$$\eta(\theta) = E_{s_0, a_0, \dots \sim \pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t) \right] \quad s_0 \sim \rho_0(s_0), \quad a_t \sim \pi(a_t | s_t), \quad s_{t+1} \sim P(s_{t+1} | s_t, a_t)$$

- Advantage Policy Gradient

$$\nabla_\theta \eta(\theta) = E_{s_0, a_0, \dots \sim \pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t | s_t) A_{\pi_\theta}(s_t, a_t) \right]$$

TRPO

- Advantage Policy Gradient

$$\nabla_{\theta} \eta(\theta) = E_{s_0, a_0, \dots \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) A_{\pi_{\theta}}(s_t, a_t) \right]$$

- 1) online learning: difficult to increase the number of data
- 2) small change in parameter may cause big change in policy

- Define a region of parameter change that can limit the change in policy
 - Trust Region
- **Theorem:** Improvement of new policy over current policy

$$\eta(\tilde{\theta}) = \eta(\theta) + E_{s_0, a_0, \dots \sim \pi_{\tilde{\theta}}} \left[\sum_{t=0}^{\infty} \gamma^t A_{\pi_{\theta}}(s_t, a_t) \right]$$

(*) Proof

Since $A_{\pi_\theta}(s_t, a_t) = r(s_t, a_t) + \gamma V_{\pi_\theta}(s_{t+1}) - V_{\pi_\theta}(s_t)$

$$\begin{aligned} & E_{s_0, a_0, \dots \sim \pi_{\tilde{\theta}}} \left[\sum_{t=0}^{\infty} \gamma^t A_{\pi_\theta}(s_t, a_t) \right] \\ &= E_{s_0, a_0, \dots \sim \pi_{\tilde{\theta}}} \left[\sum_{t=0}^{\infty} \gamma^t \left(r(s_t, a_t) + \gamma V_{\pi_\theta}(s_{t+1}) - V_{\pi_\theta}(s_t) \right) \right] \\ &= \eta(\tilde{\theta}) + E_{s_0, a_0, \dots \sim \pi_{\tilde{\theta}}} \left[\sum_{t=0}^{\infty} \gamma^{t+1} V_{\pi_\theta}(s_{t+1}) - \sum_{t=0}^{\infty} \gamma^t V_{\pi_\theta}(s_t) \right] \\ &= \eta(\tilde{\theta}) + E_{s_0, a_0, \dots \sim \pi_{\tilde{\theta}}} \left[\sum_{t=1}^{\infty} \gamma^t V_{\pi_\theta}(s_t) - \sum_{t=0}^{\infty} \gamma^t V_{\pi_\theta}(s_t) \right] \\ &= \eta(\tilde{\theta}) - E_{s_0, a_0, \dots \sim \pi_{\tilde{\theta}}} [V_{\pi_\theta}(s_0)] = \eta(\tilde{\theta}) - \eta(\theta) \end{aligned}$$

TRPO

- Due to the theorem, objective function is now

$$E_{s_0, a_0, \dots \sim \pi_{\tilde{\theta}}} \left[\sum_{t=0}^{\infty} \gamma^t A_{\pi_{\theta}}(s_t, a_t) \right]$$

- Above objective function is evaluated using data from new policy $\pi_{\tilde{\theta}}$
- Can we somehow use the data generated from previous policy π_{θ} ?
 - how to change $\pi_{\tilde{\theta}}$ to π_{θ}

Discounted Visitation Frequency

- Discounted visitation frequency: the probability of agent staying state s using policy π_θ

$$\rho_{\pi_\theta}(s) = P_{\pi_\theta}(s_0 = s) + \gamma P_{\pi_\theta}(s_1 = s) + \gamma^2 P_{\pi_\theta}(s_2 = s) + \dots = \sum_{t=0}^{\infty} \gamma^t P_{\pi_\theta}(s_t = s)$$

- Objective function is redefined using this frequency

$$J(\tilde{\theta}) = \eta(\tilde{\theta}) - \eta(\theta) = \sum_s \rho_{\pi_{\tilde{\theta}}}(s) \sum_a \pi_{\tilde{\theta}}(a|s) A_{\pi_\theta}(s, a)$$

- Can define this function using data from old policy π_θ ?
- We can assume $\rho_{\pi_{\tilde{\theta}}}(s) \approx \rho_{\pi_\theta}(s)$
- Therefore,

$$L(\tilde{\theta}) = \sum_s \rho_{\pi_\theta}(s) \sum_a \pi_{\tilde{\theta}}(a|s) A_{\pi_\theta}(s, a)$$

- Now

$$\begin{aligned} L(\tilde{\theta}) &= \sum_s \rho_{\pi_\theta}(s) \sum_a \pi_{\tilde{\theta}}(a|s) A_{\pi_\theta}(s, a) \\ &= \sum_s \rho_{\pi_\theta}(s) \sum_a \pi_\theta(a|s) \left[\frac{\pi_{\tilde{\theta}}(a|s)}{\pi_\theta(a|s)} A_{\pi_\theta}(s, a) \right] = E_{\pi_\theta} \left[\frac{\pi_{\tilde{\theta}}(a|s)}{\pi_\theta(a|s)} A_{\pi_\theta}(s, a) \right] \end{aligned}$$

(*) Importance Sampling Perspective

- We know

$$E_{X \sim P}[f(X)] = \sum P(X)f(X) = \sum Q(X) \frac{P(X)}{Q(X)} f(X) = E_{X \sim Q} \left[\frac{P(X)}{Q(X)} f(X) \right]$$

- From

$$J(\tilde{\theta}) = \sum_s \rho_{\pi_{\tilde{\theta}}}(s) \sum_a \pi_{\tilde{\theta}}(a|s) A_{\pi_{\theta}}(s, a) = E_{\pi_{\tilde{\theta}}}[A_{\pi_{\theta}}(s, a)]$$

- Substituting

$$P(X) \rightarrow \pi_{\tilde{\theta}}(s|a), Q(X) \rightarrow \pi_{\theta}(s|a), f(X) \rightarrow A_{\pi_{\theta}}(s, a)$$

- We have

$$\begin{aligned} J(\tilde{\theta}) &= E_{\pi_{\tilde{\theta}}}[A_{\pi_{\theta}}(s, a)] = E_{\pi_{\theta}} \left[\frac{\pi_{\tilde{\theta}}(a|s)}{\pi_{\theta}(a|s)} A_{\pi_{\theta}}(s, a) \right] \\ &= \sum_s \rho_{\pi_{\theta}}(s) \sum_a \pi_{\theta}(a|s) \left[\frac{\pi_{\tilde{\theta}}(a|s)}{\pi_{\theta}(a|s)} A_{\pi_{\theta}}(s, a) \right] \end{aligned}$$

Trust Region

- However $E_{\pi_\theta} \left[\frac{\pi_{\tilde{\theta}}(a|s)}{\pi_\theta(a|s)} A_{\pi_\theta}(s, a) \right]$ has high variance

- We know that $Var[Z] = E[Z^2] - E[Z]^2$

- Let $Z = \frac{P(X)}{Q(X)} f(X)$, then

$$Var[Z] = E_{X \sim Q} \left[\left(\frac{P(X)}{Q(X)} f(X) \right)^2 \right] - E_{X \sim Q} \left[\frac{P(X)}{Q(X)} f(X) \right]^2 = E_{X \sim P} \left[\frac{P(X)}{Q(X)} f(X)^2 \right] - E_{X \sim P} [f(X)]^2$$

- Let $P(X) \Rightarrow \pi_{\tilde{\theta}}(a|s)$, $Q(X) \Rightarrow \pi_\theta(a|s)$, $f(X) \Rightarrow A_{\pi_\theta}(s, a)$
- Therefore, big changes in policy causes high variance

$$\frac{P(X)}{Q(X)} = \frac{\pi_{\tilde{\theta}}(a|s)}{\pi_\theta(a|s)}$$

Trust Region

- Difference between policies should be small
- We define a Trust region that is based on difference between two policies
 - We use Kullback-Leibler(KL) divergence
- Kullback-Leibler(KL) divergence between two probabilities $p(x)$ and $q(x)$

$$D_{KL}(p||q) = E_{p(x)} \left[\log \frac{p(x)}{q(x)} \right] = \sum_x p(x) \log \frac{p(x)}{q(x)}$$

- Difference between two policies using KL divergence

$$D_{KL} \left(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s) \right) = E_{\pi_{\theta}} \left[\log \frac{\pi_{\theta}(a|s)}{\pi_{\tilde{\theta}}(a|s)} \right] = \sum_a \pi_{\theta}(a|s) \log \frac{\pi_{\theta}(a|s)}{\pi_{\tilde{\theta}}(a|s)}$$

Trust Region

- Therefore new parameter $\tilde{\theta}$ should satisfy the following condition

$$\max_s D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s)) \leq \delta$$

- This condition is too strict. Therefore we normally use average value

$$E_{\pi_{\theta}} \left[D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s)) \right] \leq \delta$$

- Finally TRPO objective function with restriction is as follows

$$\begin{aligned} & \operatorname{argmax}_{\tilde{\theta}} E_{\pi_{\theta}} \left[\frac{\pi_{\tilde{\theta}}(a | s)}{\pi_{\theta}(a | s)} A_{\pi_{\theta}}(s, a) \right] \\ & \text{subject to } E_{\pi_{\theta}} \left[D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s)) \right] \leq \delta \end{aligned}$$

Computing TRPO

- Computing $\tilde{\theta}$ is not easy
- We use approximate function for objective function($L(\tilde{\theta})$) and trust region($l(\tilde{\theta})$)

- Functions are approximated using Taylor Series:

$$\tilde{f}(x) = f(c) + \nabla f(c)^T (x - c) + \frac{1}{2!} (x - c)^T H(x - c)$$

- Let $L(\tilde{\theta}) = E_{\pi_{\theta}} \left[\frac{\pi_{\tilde{\theta}}(a|s)}{\pi_{\theta}(a|s)} A_{\pi_{\theta}}(s, a) \right]$

- Using Taylor Series, $L(\tilde{\theta})$ is approximated as follows

$$L(\tilde{\theta}) \approx L(\theta) + \nabla_{\tilde{\theta}} L(\theta)^T (\tilde{\theta} - \theta)$$

$$L(\tilde{\theta}) \approx \nabla_{\tilde{\theta}} L(\theta)^T (\tilde{\theta} - \theta)$$

Computing TRPO

- Now define approximate function for trust region

- Let $l(\tilde{\theta}) = E_{\pi_{\theta}}[D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s))] = E_{\pi_{\theta}} \left[\sum_a \pi_{\theta}(a|s) \log \frac{\pi_{\theta}(a|s)}{\pi_{\tilde{\theta}}(a|s)} \right]$

- Using Taylor Series

$$l(\tilde{\theta}) \approx l(\theta) + \nabla_{\tilde{\theta}} l(\theta)^T (\tilde{\theta} - \theta) + \frac{1}{2} (\tilde{\theta} - \theta)^T \nabla_{\tilde{\theta}}^2 l(\theta) (\tilde{\theta} - \theta)$$

- We know that

1) $l(\theta) = E_{\pi_{\theta}}[D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\theta}(\cdot | s))] = 0$

2) (*) $\nabla_{\tilde{\theta}} l(\theta) = \nabla_{\tilde{\theta}} D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s)) \Big|_{\tilde{\theta}=\theta} = \nabla_{\tilde{\theta}} \sum_a \pi_{\theta}(a|s) \log \frac{\pi_{\theta}(a|s)}{\pi_{\tilde{\theta}}(a|s)} \Big|_{\tilde{\theta}=\theta}$

$$= - \sum_a \nabla_{\tilde{\theta}} \pi_{\tilde{\theta}}(a|s) = - \nabla_{\tilde{\theta}} \sum_a \pi_{\tilde{\theta}}(a|s) = 0$$

- Due to 1) and 2), approximation function of $l(\tilde{\theta})$ is

$$l(\tilde{\theta}) \approx \frac{1}{2} (\tilde{\theta} - \theta)^T \nabla_{\tilde{\theta}}^2 l(\theta) (\tilde{\theta} - \theta)$$

Computing TRPO

- Solving in terms of $\tilde{\theta}$ (H : Hessian matrix of $D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s))$)

$$\tilde{\theta} = \theta + \sqrt{\frac{2\delta}{g^T H^{-1} g}} H^{-1} g$$

- This is called **natural policy gradient**
- Above is obtained based on the approximate function using the Taylor series.
- Accordingly, the range of variation of $\tilde{\theta}$ value may deviate from the trust region
- When $\tilde{\theta}$ value is out of trust region, the j value is increased by one so that the policy change by $\tilde{\theta}$ is within trust region ($x \approx H^{-1} g$)

$$\tilde{\theta} = \theta + \alpha^j \sqrt{\frac{2\delta}{x^T H x}} x$$

- This is called **backtracking line search**

Proximal Policy Optimization(PPO)

- TRPO needs a lot of computation
 - Computing Hessian Matrix and inverse matrix
- PPO: Approximate method of TRPO

PPO Adaptive KL Penalty

Method 1)

- TRPO problem is given as

$$\operatorname{argmax}_{\tilde{\theta}} E_{\pi_{\theta}} \left[\frac{\pi_{\tilde{\theta}}(a|s)}{\pi_{\theta}(a|s)} A_{\pi_{\theta}}(s, a) \right] \text{ subject to } E_{\pi_{\theta}} \left[D_{KL} \left(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s) \right) \right] \leq \delta$$

- Using Lagrangian dual, objective function is transformed as

$$\operatorname{argmax}_{\tilde{\theta}} E_{\pi_{\theta}} \left[\frac{\pi_{\tilde{\theta}}(a|s)}{\pi_{\theta}(a|s)} A_{\pi_{\theta}}(s, a) \right] - \beta \left(E_{\pi_{\theta}} \left[D_{KL} \left(\pi_{\theta}(\cdot | s) || \pi_{\tilde{\theta}}(\cdot | s) \right) \right] \right)$$

PPO

Method 2)

Another way of computing β

- Let

$$d = E_{\pi_{\theta}} \left[D_{KL} \left(\pi_{\theta}(\cdot | s), \pi_{\tilde{\theta}}(\cdot | s) \right) \right]$$

- Adjust β using the following

$$\beta = \begin{cases} \beta/2, & d < \frac{\delta}{1.5} \\ \beta * 2, & d > \delta * 1.5 \end{cases}$$

PPO

PPO with Clipped Objective

- If the ratio of the two policies exceeds a certain range, clip the value and forcing it to be a small value


- Clip function is defined as

$$\text{clip}(x, 1 - \epsilon, 1 + \epsilon) = \begin{cases} 1 - \epsilon, & \text{if } x < 1 - \epsilon \\ x, & \text{if } 1 - \epsilon \leq x \leq 1 + \epsilon \\ 1 + \epsilon, & \text{if } x > 1 + \epsilon \end{cases}$$

- Object function of PPO is given as

$$\mathcal{L}^{\text{CLIP}}(\tilde{\theta}) = E_{\pi_{\theta}} \left[\min \left(r_t(\tilde{\theta}), \text{clip}(r_t(\tilde{\theta}), 1 - \epsilon, 1 + \epsilon) \right) A_{\pi_{\theta}}(s_t, a_t) \right]$$

$$\text{where } r_t(\tilde{\theta}) = \frac{\pi_{\tilde{\theta}}(a_t|s_t)}{\pi_{\theta}(a_t|s_t)}$$



DDPG(Deep Deterministic Policy Gradient)

- Computing the maximum over actions in the target is impossible in continuous action spaces.
- DDPG deals with this by using a **target policy network** to compute an action which approximately maximizes
- The target policy network is found the same way as the target Q-function:
- DPG(Deterministic policy gradient): computing gradient on continuous action
- DDPG(Deep deterministic policy gradient): improved DPG using DQN techniques and others

DDPG

- Value network (DQN)



- Error function of value network is (w : parameter of value network)

$$J(w) = E_{s,a,r,s'} \left(r + \gamma \max_{a'} \hat{q}_w(s', a') - \hat{q}_w(s, a) \right)^2$$

- When using target value network in DQN (w' : parameter of target network)

$$J(w) = E_{s,a,r,s'} \left[\left(r + \gamma \max_{a'} \hat{q}_{w'}(s', a') - \hat{q}_w(s, a) \right)^2 \right]$$

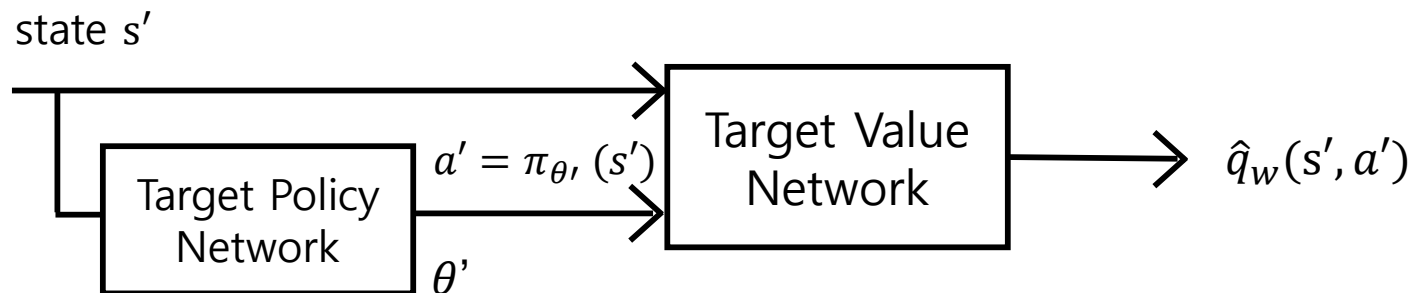
- Since actions are continuous, $\max_{a'} \hat{q}_{w'}(s', a')$ is not computable

DDPG

- In DDPG target network, finding best a' in $\max_{a'} \hat{q}_{w'}(s', a')$ is done by target policy network
- Target policy network: A new policy network that finds the optimal behavior a' from the state s'
- The target policy network is a policy network that outputs an action a' that maximizes the Q value for a particular state input s (θ' is parameter of target policy network): $a' = \pi_{\theta'}(s)$



- Now, in target value network, $\max_{a'} \hat{q}_{w'}(s', a')$ is replaced by $\hat{q}_{w'}(s', \pi_{\theta'}(s'))$



Training Value Network

- Error function of value network is defined as ($\max_{a'} \hat{q}_{w'}(s', a') \rightarrow \hat{q}_{w'}(s', \pi_{\theta'}(s'))$)

$$J(w) = E_{s,a,r,s' \sim D} \left[\left(\underbrace{r + \gamma \hat{q}_{w'}(s', \pi_{\theta'}(s'))}_{\substack{\text{target value network} \quad \text{target policy} \\ \text{network}}} - \underbrace{\hat{q}_w(s, a)}_{\text{value network}} \right)^2 \right]$$

- Differentiate $\hat{q}_w(s, a)$ only (semi-gradient)
- Update w using
$$w = w - \alpha \nabla_w J(w)$$
- Periodically update target network parameter with value network
 - Or Polyak update

Training Policy Network

- Now train policy network
- Since actions are continuous, there is only one output in policy network



- θ : parameter of policy network
 - handles deterministic policy only
- The learning of the policy network: learn the parameter θ of the policy network so that Q value of $(s, \pi_{\theta}(s))$ is maximized.
 - Therefore, objective function of policy network is

$$q_w(s, \pi_{\theta}(s))$$

- Using experience replay D

$$J(\theta) = E_D[q_w(s, \pi_{\theta}(s))]$$

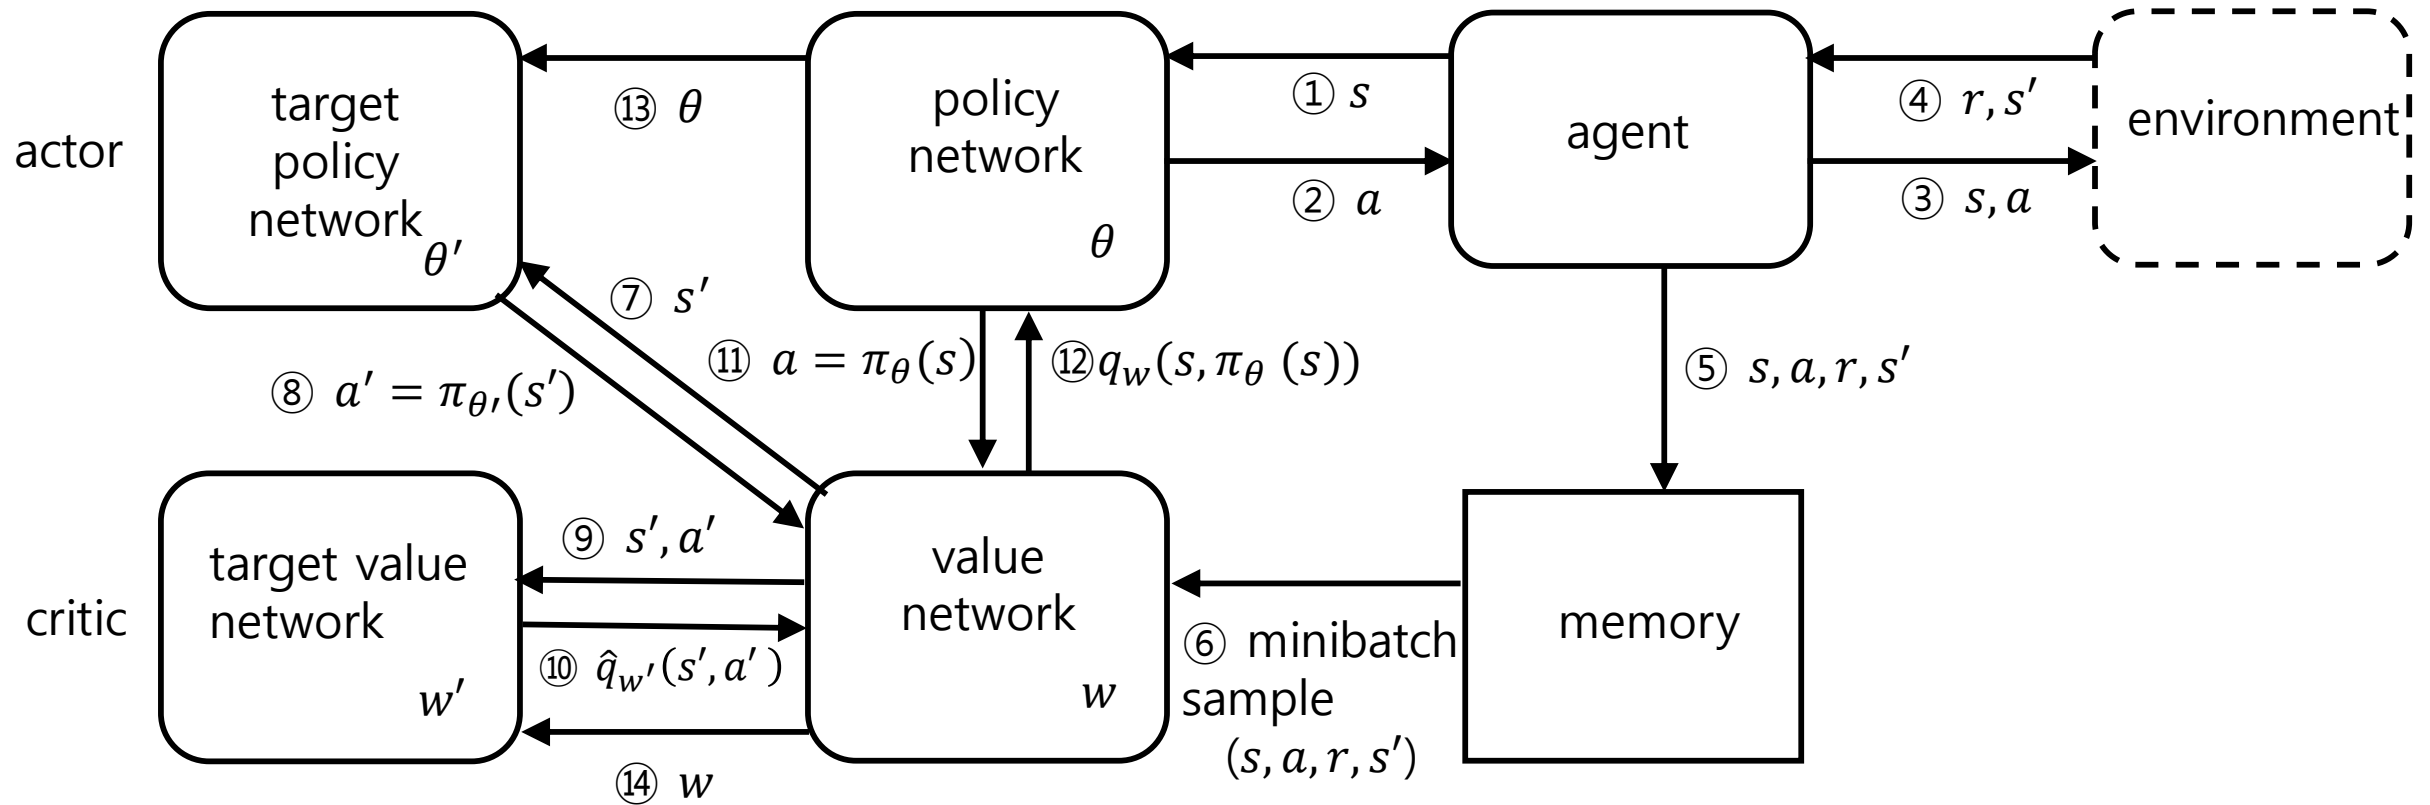
Training Policy Network

- Given objective function of policy network

$$J(\theta) = E_D[q_w(s, \pi_\theta(s))]$$
$$\nabla_\theta J(\theta) = E_D[\nabla_{\pi_\theta(s)} q_w(s, \pi_\theta(s)) \nabla_\theta \pi_\theta(s)]$$

- About $\nabla_{\pi_\theta(s)} q_w(s, \pi_\theta(s))$: differentiate value network q_w in terms of network input $\pi_\theta(s)$
- This means how much the input value a affects in learning q_w
- Differentiating networks with respect to inputs are easily implemented in deep learning libraries(e.g., Pytorch).
- When choosing action a from policy network, we use the following to increase randomness
$$a = \pi_\theta(s) + e$$
- Periodically, parameters(θ) of policy network are copied to those(θ') of target policy network
 - Polyak update

DDPG



Algorithm 1 Deep Deterministic Policy Gradient

- 1: Input: initial policy parameters θ , Q-function parameters ϕ , empty replay buffer \mathcal{D}
- 2: Set target parameters equal to main parameters $\theta_{\text{targ}} \leftarrow \theta$, $\phi_{\text{targ}} \leftarrow \phi$
- 3: **repeat**
- 4: Observe state s and select action $a = \text{clip}(\mu_{\theta}(s) + \epsilon, a_{\text{Low}}, a_{\text{High}})$, where $\epsilon \sim \mathcal{N}$
- 5: Execute a in the environment
- 6: Observe next state s' , reward r , and done signal d to indicate whether s' is terminal
- 7: Store (s, a, r, s', d) in replay buffer \mathcal{D}
- 8: If s' is terminal, reset environment state.
- 9: **if** it's time to update **then**
- 10: **for** however many updates **do**
- 11: Randomly sample a batch of transitions, $B = \{(s, a, r, s', d)\}$ from \mathcal{D}
- 12: Compute targets

$$y(r, s', d) = r + \gamma(1 - d)Q_{\phi_{\text{targ}}}(s', \mu_{\theta_{\text{targ}}}(s'))$$

- 13: Update Q-function by one step of gradient descent using

$$\nabla_{\phi} \frac{1}{|B|} \sum_{(s, a, r, s', d) \in B} (Q_{\phi}(s, a) - y(r, s', d))^2$$

- 14: Update policy by one step of gradient ascent using

$$\nabla_{\theta} \frac{1}{|B|} \sum_{s \in B} Q_{\phi}(s, \mu_{\theta}(s))$$

- 15: Update target networks with **Polyak update**

$$\begin{aligned}\phi_{\text{targ}} &\leftarrow \rho \phi_{\text{targ}} + (1 - \rho) \phi \\ \theta_{\text{targ}} &\leftarrow \rho \theta_{\text{targ}} + (1 - \rho) \theta\end{aligned}$$

- 16: **end for**
- 17: **end if**
- 18: **until** convergence

<https://spinningup.openai.com/en/latest/algorithms/ddpg.html>

TD3(Twin Delayed Deep Deterministic Policy Gradient)

- TD3 is very similar to DDPG
- But, DDPG has a few problems:
 - 1) DDPG has Q value overestimation problem
 - 2) Since DDPG learns deterministic policy, it is necessary to increase randomness in action selection

Q overestimation problem:

- In DQN, it maintains two independent networks Q_A and Q_B
- Q_A is used for evaluating and Q_B is used for actual Q value
- In TD3, we choose minimum value between Q_A and Q_B
- $\min(Q_A, Q_B)$ becomes target of Q_A and Q_B network, respectively

TD3

Randomness in action

- Since DDPG handles deterministic policy, we need to increase the degree of exploration
- DDPG also adds randomness when choosing action a (p. 40)
 - when choosing action a from policy network, it adds a bit of random value
- In TD3, it add randomness when choosing action a' from state s' from experience replay data as well

$$(s, a, r, s')$$

- Therefore, action a' is selected using
$$a' = \pi_{\theta'}(s') + \epsilon, \text{ where } \epsilon \sim \text{clip}(N(0, \sigma), -c, c)$$
 - Assume noise ϵ follows Gaussian distribution $N(0, \sigma)$
 - To preventive excess noise, we also clip the noise
- Also to prevent error from happening, action values are clipped as well.

$$a_{low} \leq a' \leq a_{high}$$

- Therefore, action a' is chosen based on
$$a' = \text{clip}(\pi_{\theta'}(s') + \text{clip}(\epsilon, -c, c), a_{low}, a_{high})$$

TD3

- Both Q-functions use a single target, calculated using whichever of the two Q-functions gives a smaller target value

$$y(r, s') = r + \gamma \min \{ Q_{w'_A}(s', a'), Q_{w'_B}(s', a') \}$$

- And then both are learned by regressing to this target

$$E_{s,a,r,s' \sim D} \left[(y(r, s') - Q_{w_A}(s, a))^2 \right]$$
$$E_{s,a,r,s' \sim D} \left[(y(r, s') - Q_{w_B}(s, a))^2 \right]$$

- The policy is learned just by maximizing

$$J(\theta) = E_{s,a,r,s' \sim D} [q_{w_A}(s, \pi_\theta(s))]$$
$$\nabla_\theta J(\theta) = E_{s,a,r,s' \sim D} [\nabla_{\pi_\theta(s)} q_{w_A}(s, \pi_\theta(s)) \nabla_\theta \pi_\theta(s)]$$

- same as DDPG
- In TD3, the policy is updated less frequently than the Q-functions are.

TD3

Initialize value network parameter w_A, w_B ; Initialize policy network parameter θ
Target value network $w'_A \leftarrow w_A, w'_B \leftarrow w_B$; target policy network parameter $\theta' \leftarrow \theta$
for $t = 1$ to T do
 choose action a in state s ($a \sim \pi_\theta(s) + \epsilon$ where $\epsilon \sim \text{clip}(N(0, \sigma), -c, c)$)
 generate reward r and next state s'
 store (s, a, r, s') in memory R
 sample minibatch from R
 for each (s_i, a_i, r_i, s'_i) in R
 $a'_i = \text{clip}(\pi_{\theta'}(s'_i) + \text{clip}(N(0, \sigma), -c, c), a_{\text{low}}, a_{\text{high}})$
 $y(r_i, s'_i) = r_i + \gamma \min \{Q_{w'_A}(s'_i, a'_i), Q_{w'_B}(s'_i, a'_i)\}$
 update value network parameter w_A and w_B , respectively
 $\underset{w}{\operatorname{argmin}} \frac{1}{N} \sum_i (y(r_i, s'_i) - Q_{w_A}(s_i, a_i))^2$ & $\underset{w}{\operatorname{argmin}} \frac{1}{N} \sum_i (y(r_i, s'_i) - Q_{w_B}(s_i, a_i))^2$
 if $(t \bmod d)$ then
 update policy network parameter θ

$$\nabla_{\theta} J(\theta) = \frac{1}{N} \sum_i \nabla_{\pi_{\theta}(s_i)} q_{w_A}(s_i, \pi_{\theta}(s_i)) \nabla_{\theta} \pi_{\theta}(s_i)$$

 update target network parameters using Polyak method
 $w'_A = \kappa w_A + (1 - \kappa) w'_A$
 $w'_B = \kappa w_B + (1 - \kappa) w'_B$
 $\theta' = \kappa \theta + (1 - \kappa) \theta'$
 end if
end for